

Practical Reasons

Many anti-realists have charged that if moral realism were true, there would be *no reason* for acting morally. In this chapter, I argue instead that only the intuitionist has a plausible account of our reasons for being moral.

7.1 The Humean argument against realism

What reasons do we have for behaving morally? There are at least two interpretations of this question, because there are two senses of ‘reason’: there are motivating reasons and normative reasons.

Motivating reasons (or ‘motives’) are the sort of things that explain why people do what they do. One’s having a motivating reason for doing *A* does not entail that one will do *A*; it means only that one has at least *some* tendency or temptation to do *A*, given the opportunity. For example, my hunger, combined with my belief that the sandwich before me is food, constitutes a motive for me to eat the sandwich.

Normative reasons are the sort of things that explain why people *should* do certain things, from the standpoint of reason.¹ A normative reason for performing an action is a consideration in favor of the action, though not necessarily a decisive consideration. Like motivating reasons, normative reasons can conflict. What it is rational to do is what one has the *most* normative reason to do. For example, the fact that I need food, together with the fact that the sandwich before me is food, constitutes a normative reason for me to eat the sandwich.

We typically have normative reasons and motivating reasons for

the same things. In the above example, I have both a motivating and a normative reason for eating the sandwich. But often the *strength* of our motives fails to match the strength of our normative reasons. Suppose I have a choice between a nutritious veggie sandwich and a delicious ice cream sandwich. I may have *better reasons*, all things considered, for eating the veggie sandwich, but I may find, alas, that I have *more motivation* to eat the ice cream. The anticipated gustatory pleasure of the ice cream is *some* justification for choosing it (insofar as pleasure is a good), but this justification may be rationally outweighed by long-term health considerations. These long-term health considerations produce in me *some* tendency to choose the veggies, but this tendency may be outweighed, motivationally, by the desire for sugary pleasure, and I may thus make the irrational choice.

The following three principles about reasons for action are widely accepted in philosophy; the first two are standard assumptions in economics as well:

- i) *The Humean Conception of Reasons* (hereafter, 'the Humean Conception'): All reasons for action derive from pre-existing desires. In particular,
 - a) You have a motivating reason to do *A* if and only if there is something you desire for its own sake that you believe *A* would or might accomplish. The strength of your reason to do *A* is proportional to the strength of that desire.
 - b) You have a normative reason to do *A* if and only if there is something you desire for its own sake that you have reason to believe *A* would or might accomplish.²
- ii) *The Belief/Desire Gap*: Beliefs don't entail desires: it is possible to believe any given objective fact obtains, while lacking any given desire.
- iii) *The Magnetism of Values*:
 - a) If one accepts a moral value, one necessarily has a motivating reason for action. For instance, if you take stealing to be wrong, then you must have at least some motivation to avoid stealing.
 - b) Moral facts, if there are any, necessarily provide normative reasons for action. If stealing is wrong, or at least if one has reason to believe it to be wrong, then one has a normative reason to avoid stealing.

These principles are commonly thought to refute moral realism. How so? There are two versions of the argument:³

First Version:

1. Motives for action require desires. (The Humean Conception.)
2. No belief about an objective fact entails the presence of any desire. (The Belief/Desire Gap.)
3. Therefore, no belief about an objective fact, by itself, entails the presence of any motive for action. (From 1, 2.)
4. Moral attitudes (the attitudes expressed by moral claims) entail the presence of motives for action. (The Magnetism of Values.)
5. Therefore, moral attitudes are not beliefs about objective facts. (From 3, 4.)

Second Version:

6. Normative reasons for action require desires. (The Humean Conception.)
7. No objective fact entails the presence of a desire. (Because objective facts are facts in the external world, apart from our desires and attitudes.)
8. No objective fact, by itself, entails the presence of any normative reason for action. (From 6, 7.)
9. Moral facts, if there are any, entail the presence of normative reasons for action. (The Magnetism of Values.)
10. Therefore, there are no objective moral facts. (From 8, 9.)

In brief, moral claims state reasons for action. Reasons for action, of either kind, require the existence of desires. But no objective fact, or belief about an objective fact, can by itself entail the existence of any desires. Therefore, moral claims do not merely report objective facts; they must contain some subjective or non-cognitive element.

There are three ways for a moral realist to respond: reject the Humean Conception, reject the Belief/Desire Gap, or reject the Magnetism of Values. Each way has been taken by at least some philosophers.⁴ I shall criticize the Humean Conception and propose an alternative conception of reasons.

7.2 The connection between motivating and normative reasons

A natural response to the first version of the Humean argument is to question (4). If I take stealing to be wrong, then perhaps it follows that I *ought* to be, or that I *think* I ought to be, motivated to avoid stealing. But it does not obviously follow that I *am* in fact motivated

to avoid stealing. I might be an immoral person, even by my own lights—that is, I might fail to be motivated by what I take to be morally correct. For similar reasons, we might doubt (6) in the second version. If I have a *normative* reason to help Jimmy with his homework, perhaps it follows that I *should* desire to help Jimmy, or that I would desire to do so *were I fully rational*, but it does not obviously follow that I *in fact* have some desire that would be satisfied by helping Jimmy. There is reason to suspect that the argument trades on an equivocation: the Humean thesis—that reasons for action require desires—is plausible only if we take ‘reasons’ to mean ‘motivating reasons’. But the Magnetism of Values premise—that moral claims state reasons for action—is plausible only if we take ‘reasons’ to mean ‘normative reasons’. That is, it is plausible that our being *motivated* to do something requires our having a desire appropriately related to the action, but not that our having a *normative* reason to do something requires such a desire; and it is plausible that our justifiedly taking something to be good must give us a *normative* reason to pursue it, but not that it must actually motivate us to pursue it. But the argument requires both the Humean Conception and the Magnetism of Values to be true for the same sense of ‘reason’.

That is not my ultimate response to the argument, however. I raise this point to show why a proponent of the Humean argument should want to defend some strong connection between motivating and normative reasons, roughly to the effect that the two kinds of reasons always go together. Fortunately for them, they can do so by means of at least two plausible arguments.

First, an argument from the ‘*ought* implies *can*’ principle: To say that you *should* do *A* implies that you *can* do *A* of your own free will.⁵ This implies that you can freely choose to do *A*, which in turn implies that you can be motivated to do *A*; a person cannot freely choose to do something while having no motivation at all for doing it. Now, to say that you have a normative reason for doing *A* implies that you *should* do *A* (provided there are no countervailing considerations). It follows that you are capable of being motivated to do *A* (provided there are no countervailing considerations). But if you had no *motive* that would appear to you to be served by doing *A*, then, since you also have no choice about what motives you have, you would be *incapable* of being motivated to do *A*. Therefore, the claim that you have a normative reason for doing *A* implies that you have some available motivation for doing *A*.⁶ In short:

1. If *S* has a normative reason for doing *A*, then *S* should do *A*, provided there are no countervailing reasons.

2. If *S* should do *A*, then *S* can freely do *A*.
3. If *S* can freely do *A*, then *S* can be motivated to do *A*.
4. Therefore, if *S* has a normative reason for doing *A*, then *S* can be motivated to do *A*, provided there are no countervailing reasons. (From 1, 2, 3.)
5. If *S* has no available motivating reason to do *A*, then *S* can not be motivated to do *A*, even if there are no countervailing reasons against *A*.
6. Therefore, if *S* has a normative reason for doing *A*, then *S* has an available motivating reason for doing *A*. (From 4, 5.)

To illustrate, suppose someone tells me that I have a normative reason to donate \$500 to Oxfam. Suppose, however, that I am so lacking in concern for my fellow human beings that I have no desire to help the poor, nor any other desire that would be satisfied by donating this money. Suppose that after thorough and correct deliberation, I realize these facts. Furthermore, I cannot help any of this; I cannot help the fact that I don't care about the poor, nor the fact that I know that donating \$500 to Oxfam satisfies none of my desires. In this situation, I cannot be motivated to donate the money. Therefore, I simply cannot donate the money of my own free will. And therefore, it is false to say that I *should* donate the money. Conversely, for it to be *true* that I should donate the money, one of the above assumptions would have to be false; I would have to have some desire that I thought donating the money might satisfy.

That is an argument a Humean might make, to draw the connection between motivating and normative reasons, and thereby to bolster the contention that '*S* should do *A*' implies that *S* has some sort of desire.

A second argument for this connection I call the Argument from Interpretation. Our normal way of interpreting other people's behavior (including their statements) relies on an assumption of rationality: when you are trying to figure out what another person believes, desires, and so on, you look at his behavior and attribute to him the beliefs, desires, and so on that would make that behavior make the most sense. For example, if someone eats a lot of chocolate, I infer that she likes chocolate and thus is undertaking the correct means of satisfying her preferences. If I instead assumed that she was irrational, I might hypothesize that she *hates* chocolate and that is why she eats it. A similar point applies to attributing beliefs to others: I assume my chocolate-eating friend believes that she is eating chocolate, because that is what it would make sense for her to believe, given her actual circumstances. If I took her to be irrational,

I might instead hypothesize that she thinks she is swimming the English Channel, and she likes swimming the Channel, so that is why she eats the chocolate. The hypotheses that assume irrationality might 'explain' the same behavior, but we rightly regard them as *prima facie* ridiculous. The point is that we need to assume others are at least generally rational in order to select an appropriate account of their behavior from among the indefinitely many possible accounts.

Something like this principle is widely accepted under the name of 'the Principle of Charity'.⁷ The principle is not that all actions or beliefs are rational, but only that we should interpret others to be by and large rational, if possible. Sometimes, people's statements and behavior provide evidence that they are irrational in certain respects. The Principle of Charity does not deny this, but it holds that rationality is a default assumption, and that we should not attribute more irrationality than necessary. The point is not merely one about how we in fact interpret others, but about what is the *correct* way of interpreting others. Notice also that the principle is not an empirical one: it is not that we have, by observing people for a long time, *discovered* that they are by and large rational. Rather, the assumption of general rationality is one that we bring to our observations, as a constraint on their interpretation. We do not directly observe others' beliefs and desires, so we would have no way of finding out whether people were by and large rational, without first having some way of inferring their mental states from their behavior. That sort of inference is just what 'interpretation' is. Since the assumption of rationality is required for interpreting behavior, the assumption of rationality cannot be tested by observing people's behavior. At most, what could be discovered empirically is that people exhibit behavior which is *capable* of a rational interpretation, but we cannot determine empirically that interpretations of people as rational are *pro tanto* better than competing interpretations of people as irrational.

The Principle of Charity is almost universally accepted. But what is its basis? It seems that the principle could be correct only if we were justified in believing, in general, that people are by and large rational. But for the reasons just discussed, this belief cannot be justified by observation. So it must be *a priori*. How can that be? Here is a natural answer: because it is a necessary truth that whenever one has a normative reason for action, one has a motivating reason for action. Suppose the contrary: suppose, that is, that one can have a normative reason to do A but no motive for doing A. If this is possible, then it seems that it would also be possible for a person to lack motives, *in general*, for doing the things he had normative reasons for doing. Such a person would be generally irrational. But if

it is possible for a person to be generally irrational, then how could we know *a priori* that people are in fact by and large rational?

The preceding argument may be weakest at its first stage, where I construed the Principle of Charity to imply that we should interpret people as being by and large rational. Perhaps instead we need only interpret people as being rational *in certain respects*, which leaves open the possibility that we might discover other broad, systematic respects in which people are not rational. However, I shall not pursue this line. I find it plausible that normative reasons have the kind of 'magnetism' the Humean argument posits. The idea is not that all beings are moved by all normative reasons. The idea is that if a being can and does *recognize* a normative reason, then the being must have at least some tendency to be moved by that reason. But there may be many beings who fail to recognize many normative reasons.

So far, I have sought to state the Humean case as strongly as possible: it seems that evaluative claims imply the existence of normative reasons for action; this, in turn, implies or strongly suggests the existence of motivating reasons for action; and that implies the existence of desires that are appropriately related to the actions in question. But no objective fact can imply the existence of a desire; therefore, evaluative claims are not objective, factual claims. Where I challenge this argument is in its assumption that all motivation stems from desire.

7.3 A rationalist conception of motivation

There are at least four kinds of motivations we are subject to:

- i) *Appetites*: Appetites include such things as hunger, thirst, and lust. They are simple, often instinctive desires, whose satisfaction is normally attended by physical (that is, sensory⁸) pleasure, and/or whose frustration is attended by physical pain or discomfort.
- ii) *Emotions*: Emotional desires, such as those occasioned by anger, fear, and love, differ from appetites in that their satisfaction or frustration is normally associated with a different (emotional) kind of enjoyment or suffering, rather than sensory pleasure or pain. In addition, emotional desires require a more sophisticated kind of cognition than appetites. To be angry at someone involves perceiving something he has done as being bad in a certain way; to fear something involves perceiving it as dangerous; to love someone involves perceiving his existence and one's relationship to him as valuable; and so on. No such sophisticated perceptions are required for the appetites. Correspondingly, emotions can

have much more sophisticated objects; one may be angry about a fact that requires quite abstract concepts to be aware of—say, the fact that a student plagiarized an assignment. Our emotions commonly have longer-term effects as well—said anger may result in one's giving the student a low grade, months later. Appetites have simpler objects and shorter-term effects.

- iii) *Prudence*: Prudence is the motivation to pursue (or avoid) a thing because one recognizes it as furthering (or setting back) one's own interests. Prudence differs from, and often opposes, appetitive and emotional desires. To return to an earlier example, my appetite for ice cream may conflict with my prudential motive to maintain good health. Likewise, my anger at someone may tempt me to hit him, while prudence bids me to restrain myself.
- iv) *Impartial Reasons*: Sometimes we are motivated to act in particular ways by our recognition of what is good, fair, honest, and so on. These reasons are impartial, in the sense that they don't appeal to our own interests or to our personal preferences. This sort of motive may conflict with any of the first three; for instance, a man may desire to cheat on his wife, but believe that this would be unfaithful (and hence bad). Or a judge may consider taking a bribe, which would promote his own interests, but believe that this would be dishonest (and hence bad).

These seem to be the main kinds of motives. There are several interesting questions about them. First, what is the object of an appetitive desire—that is, when one has an appetite, what is it that one wants? One traditional view is that it is physical pleasure or the avoidance of physical pain, and this has led some thinkers to the generalization that *all* desire boils down to the desire for pleasure and the absence of pain.⁹ This, however, is a mistake. One may of course desire pleasure, but this is not what typical appetites consist in, even though pleasure normally attends the satisfaction of those appetites. When you are hungry, what you want is food. You may *also* want gustatory pleasure, but *that* desire is not what your hunger consists in; hunger consists in the desire for food. To illustrate, suppose you are hungry. Some mad scientists stimulate your brain so that you have an incredibly realistic hallucination of eating a salmon filet, and so that at the same time you stop feeling hungry. In this case, I would say that your desire (the hunger) *went away* without being *satisfied*—you did not get what you wanted, although you may have been tricked into thinking you did. Our desires are typically focused on *the world*, not on our minds.¹⁰

Similarly, what are the objects of emotional desires? If one is

tempted to view typical appetites as desires for physical pleasure, one might be similarly tempted to treat emotional desires as desires for emotional satisfaction. Thus, one might view a man's love for his children as giving rise chiefly to a desire for the emotional satisfaction of seeing them well off. But in fact, what love gives rise to is the desire that loved ones *be* better off. That is why, for example, taking steps to ensure that one's children are provided for after one's death is a manifestation of love.

To illustrate, imagine that you are offered a chance to sell your children into slavery for \$10,000. The children will be miserable for the rest of their lives. However, as part of the deal, you will receive a hypnosis session that will cause you to forget that the children were sold into slavery and will, furthermore, induce in you a permanent delusion that the children are extremely happy and successful. You will thereby enjoy the satisfaction of *thinking* your children to be well off. Assuming that you love your children very much, would you take the deal? If you answered 'no', then your chief desire is not for emotional satisfaction.¹¹

Where does the desire for pleasure fit into my taxonomy of motivations? I think that it fits under prudence, in the admittedly broad sense in which I have defined 'prudence'. The desire to secure pleasures for oneself is a desire to promote one's interests in a certain way. We may also just have an appetite for pleasure.

This brings us to another simplistic but common psychological theory: the theory that all actions are motivated by concern for self-interest. This theory falls to much the same objection as the previous one. We are moved by a variety of different desires—desires for food, sex, power, wealth, love, fame, and so on—most of which are for things logically independent of the advancement of our interests. No *recherché* or even particularly noble examples are required to refute psychological egoism. Consider a desire, prompted by anger, to hit someone: one need not take such an action to be in any way in one's interests in order to feel the desire, and one may well take it to be very much against one's interests.¹²

These remarks undermine two particular Humean theories of reasons: the view that all actions are motivated by the desire to secure pleasure and avoid pain, and the view that all actions are motivated by the desire to promote one's own interests. But they leave intact the more general Humean framework, according to which all actions are motivated by desires of *some* kind.

Where does morality fit into the taxonomy? I take moral principles to be a species of impartial reasons. Thus, considerations of what is good, fair, honest, and so on are both moral reasons and impartial

reasons for action. There may be other impartial reasons that are not moral (perhaps the fact that a certain procedure is the most efficient way of producing pencils is an impartial reason, but not a moral reason, for one to cause the pencil factory to make pencils in that way), but I won't discuss them here.

We always have an impartial reason for doing what we know to be morally correct. But this is not to say that that reason is the only, or even the primary reason we have for acting morally. We also have many non-moral reasons for being moral. Consider, for example, a person who does what he believes is right, not merely because it is right, but because he believes God will punish wrongdoers; or a person who refrains from stealing, not merely because he recognizes theft as wrong, but because stealing makes him feel bad. These would be, respectively, prudential and emotional motives for doing what is right.

How does the notion of 'desire' map onto my taxonomy of motives? In ordinary language, appetites and emotional motives are commonly called 'desires', while prudential and impartial motives are not. Imagine a student who knows she has an exam tomorrow morning, for which she has not studied. A friend invites her to a party. The dialogue might go like this:

Friend: Hey, you wanna go to a party tonight?

Student: I really want to go, but I should study tonight.

Friend: So you're saying you'd rather study than party?

Student: Of course I don't *want* to study; I'd much rather be party-ing. But I really need to study.

The natural interpretation is that the student has described her motivations correctly: she has an emotional desire to party, no desire to study, but recognizes—that is, believes—that it would be very much in her interests to study (she 'needs' to). The latter, prudential motivation may overrule her desire.

Or imagine a judge who has been offered a bribe to acquit a criminal. The dialogue might go like this:

Judge: I won't take your bribe.

Criminal: What's the matter, don't you want \$100,000?

Judge: Of course I want \$100,000. But acquitting you would be wrong.

Criminal: So you're saying that you want to send me to prison more than you want to have \$100,000? You must really hate me.

- Judge: It doesn't matter how I feel about you. You committed a crime, and criminals should be punished, and it's my job to see to it.
- Criminal: Look, I bet my money could buy you a lot more enjoyment than whatever satisfaction you'd get from seeing me punished.
- Judge: You just don't get it.

The natural interpretation is that the judge recognizes that taking the bribe would best promote his own interests, satisfy his desires, and increase his enjoyment, but he believes it would be wrong. His moral belief overrules his personal interests.

What would the Humean say? That in the first example, the student *wants* to get a good grade on the test, more than she wants to go partying, and that in the second example, the judge *wants* the criminal to be punished more than he wants the \$100,000. Is this true? Not in any non-trivial sense. 'Desire', in English, seems often to refer to appetites and emotions, as the above examples illustrate. In that sense, the Humean's claim is false. The student may well feel more desire to go partying than to get a good grade on the test, but still recognize that it would be more in her interests to study. The judge may feel more desire to take the \$100,000 than to punish the criminal, but recognize that justice requires him to punish the criminal. Introspectively, there is simply no evidence for the Humean claim.

In another sense, however, a 'desire' is simply any motivation, of whatever kind.¹³ In this sense, it is trivial that in my examples, the student desires to study and the judge desires to refuse the bribe. On this interpretation, the core Humean thesis that motives for action require desires amounts to no more than the following tautology, with which even the staunchest rationalist will agree: Motives for action require motives for action. But surely no interesting conclusion can be inferred from this. In particular, on the present conception of desire, it can no longer be assumed that beliefs are not desires; if a belief motivates action, then it would *ipso facto* be a 'desire' in the present sense. And we should simply have to distinguish two kinds of 'desires': the kind that are constituted by beliefs, and the kind that are constituted by appetites or emotions.

I take it that the Humean does not want to trivialize his thesis in this way; most Humeans will say that we have some independent conception of 'desire', as an introspectively identifiable type of mental state, on which it is not an open question whether a desire can be a belief.¹⁴ But on introspecting, the only such mental states I

can find that might be called ‘desires’ are the appetites and emotions. And it is not debatable whether I always act on my currently strongest appetites and emotions—I know that, thankfully, I do not.

7.4 Why believe the Humean conception?

7.4.1 The intuitive appeal of the Humean conception

Why have so many philosophers embraced the Humean Conception of Reasons? Most seem simply to consider it intuitively obvious—it is just obvious that reasons for action require goals, and it is just obvious that only our desires, and not mere beliefs about the world outside of us, can provide us with goals.¹⁵ Let us first discuss this appeal to intuition, and then turn to some arguments for the Humean Conception.

I am at least as willing as the next person to consider intuitions. But I remain unpersuaded by the widespread intuition supporting the Humean Conception. This intuition is a *general* one; that is, its content is a universal generalization (‘All reasons have such-and-such form’ or ‘No belief can be inherently motivating’). I think that intuitions of this kind can often be easily undermined by particular examples of things that defy the generalization. For example, consider the principle, ‘Two solid objects cannot occupy the same space at the same time’. On its face, this seems obviously true. But now consider this case: you have a lump of clay, which we can call ‘Lumpy’. You proceed to mold the clay in such a way as to form a statue. Call the statue ‘Venus’. It seems that Lumpy still exists; it has simply been molded into a different shape. Furthermore, Lumpy and Venus now occupy the same space. Both are solid objects. And they are distinct objects, since Lumpy existed before Venus did. Therefore, we have two solid objects occupying the same space at the same time. Isn’t this a deep philosophical paradox?

No, it isn’t. We should respond to the puzzle by saying: ‘Okay, if you’re going to count cases like *that*, then two solid objects can occupy the same space at the same time. When I said this was impossible, I had in mind pairs of objects like: my hand and this table. If I had been thinking of “pairs of objects” like the statue and the mass of clay it is composed of, I would never have said that two solid objects couldn’t occupy the same space, because obviously *those* things can’. Thus, I think that the general intuition we began with—even though it was quite a firm and widely-shared intuition—has no probative force against the judgment that Lumpy and Venus are in the same place, or that they are two solid objects. The reason is that

the general intuition depended on our having certain kinds of cases, and not others, in mind in the first place.

I am not claiming that this is true of *all* general intuitions, but it is true of many of them. And in particular, I think it is true of the intuition supporting the Humean Conception. When someone puts forward the general principle, 'Beliefs cannot by themselves be motivating', I think that we can only judge the plausibility of that generalization by running through different kinds of beliefs in our minds. To be sure, the principle seems plausible when we think about *most* beliefs—for instance, the belief that my house is on fire, the belief that the weasels are mammals, the belief that 13 is prime, and so on. In fact, it seems plausible when you consider any belief about any purely natural, external fact. But now consider another class of beliefs: beliefs about *value*. For instance, the belief that stealing is *morally wrong*, the belief that happiness is *good*, and so on. If you're counting *this* sort of belief, then the principle does not seem at all obvious; indeed, it seems very *implausible*. I think the principle that only desires can motivate seems obvious only when we are ignoring certain cases—just as the principle that two solid objects can't occupy the same space only seems obvious when we are ignoring certain cases.

Consider why it seems plausible that a belief about a natural fact cannot be inherently motivating. One reason is that it is unclear what would determine *which* course of action the natural fact is supposed to favor. Thus, suppose you see someone sitting in a train station. You inform him that the train for New Brunswick is about to leave. Can this fact, by itself, motivate him to act in some way—for instance, to get up and board the train? One reason for denying that it can is that there doesn't seem to be any particular course of action that uniquely bears any special relation to the fact. Granted, the fact that the train is about to leave implies that getting up and boarding the train is required *in order to get to New Brunswick on that train*. On the other hand, refraining from getting up and boarding the train is required *in order to keep out of New Brunswick*. The mere fact that the train is about to leave does not tell us whether getting to New Brunswick is good or bad, so it doesn't pick out any particular course of action as the one to take.

But evaluative facts are obviously different. They *do* tell us what is good or bad, and there is no puzzle about which course of action an evaluative judgment supports. The judgment that one should keep promises obviously bears a special relationship to promise-keeping that it does not bear to other kinds of actions.

Now consider Hume's main argument against the Rationalist

Conception: Hume says that the function of reason is only to determine what is true and false. But actions cannot be true or false. Therefore, actions cannot be evaluated by reason, and therefore they cannot be rational or irrational.¹⁶

That is an awful argument. Consider a parallel argument: Reason can only judge truth and falsity. Numbers cannot be true or false. Therefore, reason cannot make any judgments about numbers—for instance, reason cannot tell us that seven is prime. Presumably Hume would say that this is a misapplication of his principle. Presumably, he would say that although numbers themselves cannot be true or false, *propositions about them*, including the proposition that seven is prime, *can* be true or false; and that is why the judgment that seven is prime can fall under the purview of reason. But if that's right, then a parallel response applies to Hume's argument: of course actions themselves cannot be true or false, but *propositions about them*, including the proposition that some particular action is right, can be; and that is why the judgment that an action is right can fall under the purview of reason.

*7.4.2 Smith's argument

Michael Smith provides an argument similar in spirit to Hume's, but stronger. Smith's argument begins with a technical notion, that of 'direction of fit'. Here is a story to illustrate that notion. There are two men in a supermarket. One has a shopping list and is going through the supermarket collecting items from the list. The other is a detective who is tracking the movements of the first man. The detective makes a list of the things the shopper buys. By the time the shopper gets to the checkout stand, both men have identical lists in their hands. Furthermore, both of them want the items on their list to match what the shopper has in his cart. But there is a difference: the shopper is supposed to make the *world* match his *list*; the detective is supposed to make his *list* match the *world*. Thus, if the detective should notice that his list says 'Twinkies' but the shopper has no Twinkies in the cart, he can fix the problem by simply crossing out the word 'Twinkies'. But if the *shopper* should notice that his list says 'Twinkies' but he has no Twinkies in the cart, he cannot fix the problem by crossing out 'Twinkies'; rather, he has to go back and put some Twinkies in the cart. We might express this difference by saying that the shopper's list has world-to-list direction of fit; the detective's list has list-to-world direction of fit.¹⁷

Now, there is something analogous to be said about the difference between beliefs and desires: Beliefs, we might say, have 'mind-to-world direction of fit', meaning that beliefs aim to match reality. If

your belief fails to correspond to reality (that is, it is false), then the belief is thereby defective and should be changed. Desires, in contrast, have 'world-to-mind direction of fit', meaning that we aim to make reality match our desires. If a desire fails to correspond to reality (that is, it is unsatisfied), the desire is not thereby defective; rather, one has reason to try to change *the world*.¹⁸

With this background, we can state Smith's argument against the Rationalist Conception of motivating reasons:

1. Having a motivating reason requires having a goal.
2. Having a goal is being in a state with world-to-mind direction of fit.
3. Beliefs have mind-to-world direction of fit, not world-to-mind.
4. Therefore, beliefs cannot be motivating reasons.¹⁹

(4) conflicts with the rationalist conception, since rationalists say that a moral belief can constitute a motivating reason.

Where might a rationalist challenge this argument? I think the crucial premise is (2).²⁰ First, let us clarify (2) by replacing the metaphorical expression 'world-to-mind direction of fit' with its literal meaning. I take it that the literal meaning of (2) is roughly this:

Having a goal is being in a state such that, if one notices that the world does not match that state, then one will have a tendency to try to change the world to remove the discrepancy.²¹

But this remains ambiguous. What does it mean for a mental state to 'match' or 'fit' the world? This might mean either that the *content* of the mental state is realized, or that the *goal* that the mental state represents is realized. This gives rise to the following two readings of (2):

- 2_C. Having a goal is being in a state such that, if one notices that the *content* of that state is not realized, then one will have a tendency to try to change the world to bring about the content of that state.
- 2_G. Having a goal is being in a state such that, if one notices that the *goal* represented by that state is not realized, then one will have a tendency to try to change the world to bring about that goal.

To illustrate the distinction, suppose I have the goal of increasing my happiness. According to the Rationalist Conception, my having this goal might consist simply in my believing that my happiness is

good.²² Now, the *content* of this belief—the thing I believe to be the case—is: *my happiness is good*. But the *goal* that the mental state represents is simply *my happiness*. This is just part of the rationalist view—the rationalist view is that one's goal is the thing that one believes to be good. Thus, (2_c) implies: if I notice that happiness is not in fact good, then I will have a tendency to try to bring it about that happiness is good. (2_G), on the other hand, implies that, if I notice that I am not in fact happy, then I will try to bring it about that I am happy.

Obviously, if (2) means (2_c), then the rationalist should reject the premise. If I somehow discover that happiness is not good, I will not try to bring it about that happiness is good (how could I do that? I cannot change the facts of ethics). On the other hand, if (2) means (2_G), we should accept it, but then we should reject premise (3). For, to avoid equivocation, (3) will now have to be read:

- 3_G. If one notices that the goal represented by a belief is not realized, then one will have a tendency to modify one's mind, *not* to try to change the world to bring about the goal.

This is false. If I notice that my happiness (the goal represented by my 'my happiness is good' belief) is not realized, then I *will* have a tendency to try to change the world to bring about that goal, that is, to bring it about that I am happy.

Why does Smith not notice the ambiguity in (2)? Perhaps because, on *his* view of motivation, there is no such ambiguity. On the Humean view, when one has a goal, the content of one's mental state is the same thing as the goal that it represents. For instance, when I want to take over the world, the content of the desire is 'that I take over the world', and the goal it represents is also *that I take over the world*. So, on the Humean view, there is no distinction to be drawn between (2_c) and (2_G). But there *is* a distinction on the rationalist view. When one has a moral belief, if the goal that is represented is *G*, then the content of the belief is not *G*, but rather '*G* is good'.

Note that I am not, in this section, trying to refute the Humean Conception. I am only trying to show that Smith does not succeed in refuting the Rationalist Conception. For Smith to succeed, he cannot assume the Humean Conception to start with. Thus, he cannot assume that (2_c) and (2_G) are equivalent. So he must tell us which of those he means. If we are taking the Rationalist Conception seriously, there is no reason why we would accept both (2) and (3) of Smith's argument: depending on how the notion of a mental state's 'matching' reality is read, we will either find (2) obviously implausible

(if it is read as [2_C]) or find (3) obviously implausible (if it is read as [3_C]). Smith therefore fails to provide a non-question-begging argument against the Rationalist.

So far, I have criticized the motivations for the Humean Conception. In the next two sections, I propose to demonstrate its deep implausibility.

7.5 Extending the Humean conception?

David Hume himself held a very simple version of the Humean Conception. In his view, it is impossible for any action or desire to be irrational; only beliefs can be evaluated by reason:

Actions may be laudable or blameable, but they cannot be reasonable or unreasonable.²³

'Tis not contrary to reason to prefer the destruction of the whole world to the scratching of my finger. 'Tis not contrary to reason for me to chuse my total ruin, to prevent the least uneasiness of an Indian or person wholly unknown to me. 'Tis as little contrary to reason to prefer even my own acknowledg'd lesser good to my greater.²⁴

This is absurd. Even modern-day Humeans do not follow Hume this far. Examples of irrational actions are easy to come by; indeed, the strikingness that has made the second quotation above famous stems precisely from the fact that the three examples Hume gives of things 'not contrary to reason' would normally be considered *paradigms* of irrationality. The only justification Hume offers for his counter-intuitive assessments of those examples consists in a deduction from the Humean Conception of reasons. But why should we accept that conception? If a proposed theory of practical rationality is unable to account for seemingly central examples of practical irrationality—and especially if it is so anemic as to be incapable of drawing *any* distinction at all between rational and irrational behavior—then it is difficult to see why we should take that theory seriously.

Fortunately, modern-day Humeans are eager to assure us, the Humean conception entails no such results. Of course there are such things as reasonable and unreasonable decisions, and what it is rational to do is not simply determined by whatever happens to be the agent's strongest desire at the time of the choice. Rather, there is a process of rational deliberation, whose outcome determines what is the rational course of action.²⁵ This view is still broadly Humean,

because the process of rational deliberation is not taken to generate any new goals independently of the subject's desires. Hereafter, I shall refer to any view that accepts roughly the Humean Conception while adding some constraints on rational action that go beyond 'choose whatever you most want now', the 'Neo-Humean view'. I shall consider five conditions on normative reasons that have been proposed by Neo-Humeans:

- i) *The Foresight Constraint*: What normative reasons we have, as well as their strength, is determined not only by our current desires, but also by the desires we know we will have in the future.²⁶ It is thus irrational to perform an action to satisfy a current desire, if one knows that more and stronger desires will consequently be frustrated in the future. This is why the student in the example from section 7.3 would be irrational to spend the night partying.
- ii) *The Imagination Constraint*: What we have normative reason to do is partly determined by what we *would* desire, if we were to imagine relevant situations vividly and accurately.²⁷ For instance, suppose that, were I to vividly imagine the plight of the starving Africans, I would desire to help them by sending money to Oxfam; then I presently have a normative reason to help them.
- iii) *The Consistency Constraint*: If one has an inconsistent set of desires—that is, desires that logically could not all be satisfied—then one has a normative reason to change one's desires to render them consistent. For instance, if I find that I want to shoot up heroin every day, but I also want to not be an habitual drug user, then I have reason to attempt to eradicate one of those desires.
- iv) *The Coherence Constraint*: We have a normative reason to modify our desires in such a way as to yield a more coherent set of desires, just as we have reason to modify our beliefs in such a way as to yield a more coherent belief system, where coherence is a matter of how systematic and unified a collection of desires/beliefs is, the degree to which some desires/beliefs can explain others, and the like. It is thus rational to adopt new desires that would help explain and systematize our existing desires, and to relinquish desires that don't fit well with the rest of our desires.²⁸ For example, suppose you desire the well-being of human beings except for members of one specific race, and you have no explanation for your discounting of that race. Then you could achieve a more coherent overall set of desires by giving up your racial prejudice and moving to a desire for the welfare of all human beings.
- v) *The Deliberation Constraint*: We have a normative reason, before

acting, to engage in rational deliberation about what will best satisfy our desires, and we have a normative reason for doing the things that such deliberation *would* find to be most suited to satisfying our desires. A person thus acts irrationally if he does what appears to him best suited to satisfying his desires, when further deliberation would have shown some other course of action to be better.

The Neo-Humean view, I shall argue, is unstable. Once we adopt a broadly Humean conception of reasons, none of these rationality constraints can be defended. A Humean is thus left with the view, either that there is no such thing as rational or irrational action, or that it is rational merely to do whatever one most wants to do at the moment.

7.5.1 The Foresight Constraint

Suppose that you presently have no desire that *A* occur but instead desire that *A* not occur; however, you know that at some future time, your psychological set will change in such a way that you will come to desire *A* (but not as a result of any change in beliefs). Do you *now* have a reason—either motivating or normative—to take steps to bring about *A*? A Humean should answer ‘no’. Only an individual’s desires at the time of decision can help explain his actions; a desire that you will have in the future but *do not now have* cannot motivate any present actions, any more than can a desire that you used to have but no longer have. If we accept the arguments of section 7.2 for a strong connection between motivating and normative reasons, this suggests that such a future desire also gives you no *normative* reasons, any more than a merely past desire does. Compare an argument of Bernard Williams’: Williams contends that in order for something to be a normative reason for *S* to do *A*, it must be something that *could* be *S*’s actual reason for doing *A*; but only something that includes one or more of *S*’s desires could possibly be *S*’s reason for doing anything; therefore, normative reasons must involve desires.²⁹ To this we might add that only something including one or more of *S*’s *present* desires could possibly be *S*’s reason for doing anything now; therefore, the same argument leads to the conclusion that normative reasons must involve present desires.

The story of Odysseus and the Sirens illustrates this point: Odysseus plans to sail past the Sirens, whose songs have lured many sailors to their deaths. He now strongly desires that the same *not* happen to him. But Odysseus knows that once he hears the Sirens’ song, he will be overcome by a desire to sail towards them, come

what may. This desire, let us suppose, will be so strong that he will consider death a price worth paying to be closer to the Sirens. But, to repeat, he *now* desires that he *not* go to the Sirens. His solution: he has himself tied to the mast while sailing past the Sirens. Is Odysseus rational in doing this? Does the desire that he will have when he hears the Sirens' song provide him with a normative reason now for taking steps to ensure that he will be able to join the Sirens? I take it that the answers are clearly 'yes' and 'no'. This fits with what I have said the Humean should say.

Here is another illustration of the same point. Sunflower is presently a socialist college student. To her dismay, she has observed that many former socialists have 'sold out' as they became older, changing their values to conform to those of their society. Suppose Sunflower justifiably believes that when she gets older, her own personality and values will also change in a way that will lead her, without acquiring any false non-evaluative beliefs, to endorse capitalism. But, to repeat, she presently opposes capitalism and supports socialism. Does Sunflower now have a reason to work towards the maintenance of the capitalist order? I take it that the correct answer, and the answer the Humean will embrace, is 'no'.

Next, suppose that you presently have a desire of a certain strength for *A* to occur, but you know that at some future time, your desire will be stronger or weaker than it presently is. *How much* reason do you now have for trying to bring about *A*? It is obvious that how strongly I am motivated to do something must be determined by how strong my desires are now, not how strong they were or will be at some other time. And given the arguments of section 7.2, this suggests that the same is true of our normative reasons. Moreover, this answer seems to be required by what we have already said: once we accept that *what* we now have reason to do is determined by what we presently desire, we should surely also say that *how much* reason we now have for doing a thing is determined by the present strength of our desires.

But now we have surrendered the Foresight Constraint entirely. Only present desires are relevant either to what reasons we have or to how strong those reasons are. Future desires are irrelevant. Or rather, future desires can only be relevant in the way that any other fact about the world can be relevant. If I happen to have a current desire that my future desires be satisfied, then the fact that in the future I will desire *x* is relevant to my action—but only in the way that, if I happen to desire that Zbigniew Brzezinski's desires be satisfied, then the fact that Zbigniew Brzezinski desires *x* is relevant to my action. In both cases, it is my present desire that provides the motivation for

my action, and in both cases, I could have lacked that desire and still been perfectly rational. Just as Humeans say there is no rational reason why I must want someone else's desires to be satisfied, so they should say there is no rational reason why I must want my own future desires to be satisfied.³⁰

Thus, there is no apparent reason why the student in my earlier example would be irrational to spend the night partying. If she chooses to do so, it will be because her desire to party was, at the time she made the choice, stronger than her desire to pass the class, graduate, and so on (how else could her action even be understood?). Though she knew that *later* her desire to pass the class would be much stronger, this is irrelevant. Even if the student had a desire that her future desires be satisfied, this would not render her action irrational; the fact that she chose to party would merely show that, at the time she made the choice, her desire to party was stronger.

7.5.2 The Imagination Constraint

The Imagination Constraint holds that our reasons for action are partly a function of desires that we do not have but would have if we were to vividly and accurately imagine certain situations. Vivid imagination might alter our desires by causing us to notice features of a situation that we would otherwise overlook, or by enabling us to better grasp a situation that we otherwise would not understand; we might then have new conative reactions to the situation. In these and other ways, vivid imagination might cause us to experience different emotions or appetites—for example, by awakening a sense of sympathy in us (we are not here concerned with purely cognitive effects of imagination).

It is not clear, however, why a desire that I *don't have* but would have if I were to go through some imaginative exercise should be thought relevant to my actual practical reasons. Such a merely potential desire cannot serve as a motivating reason; only actual desires can move me to act. Given the connection between motivating and normative reasons, it seems that it cannot serve as a normative reason either. Granted, vividly imagining a situation may change my emotional state and so change my desires, but what makes the new set of desires any better, or 'more rational', than the old set? Perhaps if correcting some errors in my reasoning would alter my desires, there would be an intelligible sense in which the new desires would be more rational—they would be more a product of the correct use of reason. But this is not the sort of change we are discussing here.

Consider an example: suppose my doctor has told me that I need

surgery to remove my appendix. I have just arrived at the hospital on the day my surgery is scheduled for. I know that, because of my squeamish disposition, if I were to imagine the surgical procedure in vivid detail, I would be so horrified at the images of the surgeon cutting me open and so on, that I would experience an overpowering urge to flee the hospital. However, I do not presently have such an urge, and I presently desire that I go through with the surgery. What is it rational for me to do? It seems clear that I should go through with the surgery, and just refrain from vividly imagining it. In this case, the desire that I would have were I to vividly imagine the surgery provides no reason for action on my part.

But now return to the example of Oxfam and the starving Africans. Suppose I presently do not want to give money to Oxfam, because I would get much more enjoyment from spending the money on a new stereo. But I also know that if I were to vividly imagine the situation in Africa, I would feel so bad about it that I would wind up sending a lot of money to Oxfam, instead of enjoying the new stereo. On a Humean conception of reasons, why should I not just conclude, 'Well, then I better *not* vividly imagine the situation in Africa', and proceed with my stereo-purchasing plans?

7.5.3 The Consistency and Coherence Constraints

At least the Humean can defend a rational requirement of consistency and coherence for our desires, in the same way that reason requires consistent and coherent beliefs, can't he?

No, he can't. First, let us review the basic difference between beliefs and desires. Recall that Michael Smith, the very philosopher who wants to defend a coherence requirement for desires, characterizes the difference between beliefs and desires in terms of 'direction of fit'. Beliefs have 'mind-to-world direction of fit', meaning that they aim to match reality. Desires have 'world-to-mind direction of fit', meaning that we aim to make reality match our desires.

Second, let us review what is bad about inconsistency in beliefs. An inconsistent belief is one that cannot be true. Since beliefs have mind-to-world direction of fit, this marks out such beliefs as inherently defective; there is always a reason for giving up such beliefs. If a (small) set of beliefs is inconsistent, one should change one's beliefs to eliminate the inconsistency, for the same reason, namely, to attain a more accurate belief system.³¹ But this line of thinking does not apply to desires. If a desire is inconsistent, then it cannot be satisfied. But since desires have world-to-mind direction of fit, a desire's unsatisfiability does *not* mark it out as defective. Since, that is, we try (or have reason to try) to make reality conform to our desires *rather*

than to make our desires conform to reality, the fact that a desire does not, or even can not match reality is not a reason for changing the desire. Nor, therefore, does the impossibility of satisfying all the members of a *set* of desires render that set of desires defective.

To illustrate, suppose that I both *believe* God exists and *want* God to exist. But an atheist gives me an argument proving that God cannot exist.³² At this point, my belief is shown to have been seriously defective (that is, wrong), and if I am rational, I must stop having that belief. But my *desire* is not thereby shown to have been 'wrong' in any sense, nor am I irrational if I continue to wish there were a God, nor is there anything defective about a value system according to which it would be good if there were a God.

We should distinguish this point from some superficially similar but false claims. As I have said, the unsatisfiability of a *desire* constitutes no defect in the desire. However, I do not deny that the impossibility of carrying out an *intention*, realizing a *hope*, or succeeding in an *attempt* constitutes a rational reason for giving up the intention, hope, or attempt. Why? Because intentions, hopes, and attempts all involve more than just desires. An intention to do *A* involves a *belief* that one will do *A* (contrast 'I intend to win this race' with 'I want to win this race'). A hope that *p* involves a belief that *p* is possible. And an attempt, besides perhaps involving the same sort of belief as an intention, normally requires, as part of its practical justification, the belief that the attempt will or might help to further some goal; but unsuccessful attempts generally do not further one's goals. This is why intentions, hopes, and attempts are open to criticism in terms of the impossibility of their objects, whereas desires are not.

Still, one might well feel that at least some desires are irrational because their objects are impossible. For instance, it seems that I would be irrational to want 2 to be equal to 3, or to desire that it both rain and not rain at the same time. I think that if these desires are irrational, it is not merely because they cannot be satisfied; perhaps they are irrational because it is unintelligible how their objects could be *good*, or because one cannot coherently conceive of their supposed objects, making it impossible for a fully rational being to entertain those desires. But in any event, these cases scarcely concern us, since such unintelligible desires are extremely rare, if they exist at all. The issue in this section is whether the Humean can provide enough constraints to give us a plausible account of rational action. The ability to criticize these unintelligible desires would still leave the Humean with a very limited conception of normative reasons, taking us very little distance towards a robust notion of practical rationality.

I should address one argument for the claim that we have a normative reason to eliminate inconsistent desires: We want it to be the case that our desires, overall, are well-satisfied, for the more of our desires are satisfied, the happier we will be, and we want to be happy. If we have inconsistent desires, we know they won't be satisfied. Therefore, we have a good reason to try to change our desires to ones that are consistent. This is a straightforward piece of means-end reasoning: we should eliminate inconsistent desires, because doing so will help achieve an end we desire. Notice that if this argument were correct, it would apply not only to desires that cannot be satisfied, but also to desires that are *unlikely* to be satisfied or that are *difficult* to satisfy. It would show that we have a normative reason to convert to a set of easily satisfiable desires, such as the desire that squares have four sides, that rocks be made of atoms, and so on. Since world hunger is unlikely to be eliminated in the foreseeable future, we have a reason to eradicate, if possible, the desire that people not starve. Of course, we might find it difficult or impossible to eradicate that desire, but then the same might be true of any other desire. For the Humean, there is no principled difference between inconsistent desires and hard-to-satisfy desires on this score.

These last results are quite implausible; we do not think of the desire for an end to world hunger as to any extent irrational, and we do not think that a being who desired only the unavoidable would thereby be more rational than we.³³ There is a straightforward, non-Humean explanation for this: since many avoidable states of affairs are good, we see desires for them as *appropriate*. We likewise see indifference towards bad states of affairs as inappropriate. Thus, we think that the nature of world hunger is such that it calls out for some negative emotional reaction, and that a person who lacked any such reaction would thereby be morally deficient.

A Humean might attempt to accommodate this idea by treating it as a matter of further desires that we have—in particular, desires that we keep our present desires. But if this works, then the same might hold of an agent with inconsistent desires: he too may desire that he keep his present desires. Again, the Humean has no theoretical rationale for treating inconsistent desires differently from merely hard-to-satisfy desires.

Once we see this, Michael Smith's idea that we should seek a coherent set of desires loses all plausibility. The reason why epistemologists consider a coherent set of *beliefs* desirable is that such a set is more likely to be by and large true.³⁴ But since it is not the job of desires to conform to reality, one cannot sustain the parallel argument that a systematic and unified set of desires is to be preferred

on the grounds that it is more likely to be by and large satisfied. The point that desires have world-to-mind direction of fit frustrates the parallel that Smith wants to draw.

7.5.4 The Deliberation Constraint

I have just attacked some of the main tenets of the Neo-Humean account of practical rationality. What remains is perhaps the most central and putatively impregnable tenet of the view: that we at least have a normative reason for figuring out what would satisfy our current desires, and adhering to the results of such reasoning. But I shall argue that even this cannot serve as a constraint enabling some behavior to be rationally criticized.

If we accept that only current, actual desires can feature in reasons (either motivating or normative) for action, then it is plausible that we should say the same about beliefs: only a person's current, actual beliefs—not beliefs he will have at some future time, and not beliefs he would have in some counterfactual situation—can figure in the reasons for action he presently has. The arguments seem to be the same: a belief that one *does not now have* cannot have any influence on one's current behavior. Furthermore, given this, it seems that such a (non-)belief cannot feature in any normative reasons that one presently has.

To illustrate, suppose you are considering whether to enter a new career as a trapeze artist in the circus, or to stick with your job as a professor. You have plenty of time to decide. After thinking about it for five minutes, it appears to you that becoming a trapeze artist would best satisfy your desires, and you decide to do that. But suppose that, if you had only deliberated a little longer and a little more carefully, you *would have* rationally and correctly come to the conclusion that remaining a professor would better satisfy your desires. Can you be rationally criticized, on a Humean view, for joining the circus? Let us divide this into two questions:

First, *given* your beliefs and desires at the time of the decision, can you be rationally criticized for joining the circus? No. You decided to join the circus because you had a certain set of desires, and you thought that decision would best satisfy them. If this were not so—if you had expected the other option to better satisfy your desires—then there would simply be no explanation, given a Humean conception of motivation, of how you could wind up choosing the circus. The very explicability of your action guarantees its rationality in terms of your beliefs and desires at the time. The only way for a Humean to rationally criticize your action, then, is to level some criticism at your beliefs at the time of the decision.

Second, then, can you be rationally criticized for believing as you did? The Humean wants to say 'yes'. Consider two forms the criticism might take: (a) You might be criticized for something over which you had no control—for instance, for not being smarter, or not having some information that you had no reasonable way of obtaining. But this hardly seems acceptable. The sort of criticism we are interested in in the first place is an indictment of the practical rationality of your decision to become a trapeze artist—that is, a criticism to the effect that you rationally *ought* to have done otherwise than you did. A criticism of something that you had no control over, whether or not that thing led up to your action, hardly seems relevant to showing that you ought to have acted differently than you did. (b) You might be criticized for something over which you *had* control—for instance, for not spending more time deliberating or expending more effort to be careful. But this criticism will fail for the reasons given in the preceding paragraph. Your act, say, of spending a certain amount of time deliberating will itself have to be assessed as rational *given* the beliefs and desires you had at the time you performed this action, if the action is explicable at all. If you had thought that deliberating longer would best serve your ends, then (given the Humean account of motivation) you would have done so. If you did not think that deliberating longer would best serve your ends, then it was rational, given that, for you not to deliberate longer. Thus, the Humean can criticize you for not deliberating longer only if he can criticize you for some *further* belief(s). A regress looms, in which at each stage the Humean's attempted criticism of you demands some further belief for which you are criticizable. In sum, for any person, *S*, and any occurrence or state of affairs, *x*:

1. If *S* does not control *x*, then *S* cannot be rationally criticized for *x*.
2. If *S* controls *x*, then *S* can be rationally criticized for *x*, on Humean grounds, only if there is something *else* that *S* can be rationally criticized for.
3. Therefore, *S* can be rationally criticized for something, on Humean grounds, only if there is an infinite series of things *S* can be rationally criticized for.
4. There never is such an infinite series.
5. So no person can be rationally criticized, on Humean grounds, for anything.

Premise (1) is true because a person cannot be held unreasonable for failing to do something he could not do. (2) is true because, on the Humean view, an action or choice always makes sense *given* the

agent's beliefs and desires; it can only be irrational if one of those beliefs or desires is irrational. Step (3) follows (assuming we disallow circular reasoning in criticisms of actions), because the 'something else' mentioned in (2) must (by (1)) be another thing over which *S* has control. But for *S* to be criticized for this second thing, there will have to be a *third* thing over which *S* has control and for which *S* can be rationally criticized . . . and so on. I take (4) as an intuitively obvious premise. (5) follows from (3) and (4).

One might object that either (1) or (2) is false, because a person can be rationally criticized for certain beliefs (say, beliefs for which there is no evidence), without there being anything else he is criticizable for. But what sort of criticism is it that can be leveled at beliefs? One sort of criticism of a belief would be that the belief is not true, or not likely to be true.³⁵ Another sort would be a practical criticism: that one *ought not* to have adopted the belief, that one had good normative reasons not to adopt it. But the first kind of criticism appears irrelevant for our purposes. If a person has done everything that it made sense for him to do given his beliefs, *and* he could not have done anything to get better beliefs, then I do not see how any criticism of his beliefs could show that he nevertheless *ought*, rationally, to have *acted* differently. Now consider the second kind of criticism. One might hold (though most philosophers doubt this) that we can control our beliefs, and we have normative practical reasons for adopting epistemically rational beliefs.³⁶ But this would return us to the above-mentioned regress: your act of adopting an epistemically irrational belief would itself be subject to the Humean account of practical rationality, in which case any criticism of this act would require that there be something *else* for which you are criticizable.

7.5.5 Coda: the need for evaluative facts

I have argued that none of the proposed rationality constraints can be motivated within a broadly Humean conception. Furthermore, in each case we can adduce clear examples, both of rational actions or desires that violate the constraint, and of irrational ones. Therefore, some extra structure not contained in the Humean Conception is needed in a theory of practical rationality, to accommodate paradigmatic examples of rational and irrational action. That extra structure, I contend, is provided by *evaluative facts*.

Consider again the Foresight Constraint. Odysseus is rational in violating the constraint when he attaches no weight to his future desire to sail closer to the Sirens. But the partying student is irrational in violating the constraint, if she fails to attach proper weight to her

future desire to pass her exam. To a Humean, this is paradoxical: sometimes we take future desires to be relevant to what is rational, and other times we do not. If basic desires cannot be evaluated, then it seems that we should either let in all basic future desires, or exclude all basic future desires from our reasons for action. From a common sense standpoint, the solution to the puzzle is obvious: Odysseus rationally judges his future desire to be an *inappropriate* one, a desire for what is in reality bad. But the partying student rationally judges, or ought to judge, her future desires to be *appropriate* ones, desires for what is in reality good. That is why it is rational for Odysseus to intentionally frustrate his future desire, but irrational for the student to do likewise. Now, the Humean might wish to argue that Odysseus' anticipated desire would be misguided on the basis of some *other* Neo-Humean rationality constraint. But this move is unavailable if my arguments in regard to the other proposed constraints have succeeded.

Second, consider the Imagination Constraint. In the surgery case, I rationally violate this constraint by attaching no weight to the desire I would form were I to vividly imagine my upcoming surgery. But in the Oxfam case, it seems that I irrationally violate the constraint when I attach no weight to the desire I would form were I to vividly imagine the plight of the starving Africans. Again, this is a puzzle for Humeans, but the common sense explanation is straightforward: the counterfactual desire in the surgery case would be a misguided desire, a desire for what is bad. The counterfactual desire in the Oxfam case would be a fitting desire, a desire for what is good. That is why we think it rational to merely evade the former would-be desire by refraining from thinking too much about the surgery, but we do not think it rational to evade the latter desire by refraining from thinking too much about world hunger.

Third, consider the Consistency Constraint. On the assumption that God cannot exist, my desire for God to exist would rationally violate this, or a closely related constraint. On the other hand, a desire for two to equal three would seem to irrationally violate the constraint. Why? Because even if we suppose the existence of God to be impossible, we can still understand how the existence of a perfect being would be good. But we cannot understand how two's equaling three would be good.

Finally, recall Michael Smith's Coherence Constraint. If my aesthetic preferences are eclectic, I am not thereby irrational to any extent. On the other hand, it seems that the racist who has concern for all humans except the members of a specific race *is* thereby irrational. Evaluative facts once again aid in explaining this: my

aesthetic tastes need not reflect any moral facts. If I like one piece of music but dislike another, very similar one, I need not hold the view that the one piece has an objective evaluative (or even aesthetic) property that the other lacks. But one's regard for people's welfare, we think, ought to reflect, at least in part, the *worth* of human beings. Since we can find no morally relevant distinction between one race and another, we think that people of all races have equal worth as human beings. A lack of regard for members of one race therefore evinces a blindness to that worth. Cases in which the coherence constraint seems appropriate can be understood as cases in which we think there are evaluative facts that constrain appropriate desires.

In assessing the promise of this approach, I commend two questions to the reader. First, are the practical rationality assessments I have made about the preceding examples plausible? Second, is it plausible that our tendency to accept those assessments is explained by our acceptance of evaluative judgments about the appropriateness of desires, of the sort I have suggested? If so, then the standard, Humean view of the relation between desires and evaluations gets things backwards. The Humean view is that it is rational to act in accordance with our evaluative judgments only because doing so helps satisfy our desires. I maintain that it is rational to act in accordance with our desires only when we would be justified in judging the satisfaction of those desires to be good. Evaluative judgments give normative force to desires, not the other way around.

This conclusion clearly points us towards ethical intuitionism. For if there are facts that provide normative reasons for action, it seems that they must fall into one of three categories: (a) facts about our psychology; (b) facts about the external, natural world; or (c) some species of non-natural, objective facts. My criticisms of the Humean Conception show that (a) is insufficient to account for all our normative reasons, given that the only psychological facts that are plausibly taken to supply reasons for action are facts about desires. Furthermore, it is generally agreed that facts of type (b) do not by themselves provide reasons for action (see section 7.4). So only (c) is left: some of our reasons for action consist of non-natural, objective facts, such as the facts that certain things are *good* or *bad*. But only intuitionists can embrace this answer, for only intuitionists are metaethical dualists (see section 1.4): it is only the intuitionists who believe that there *are* any facts of kind (c). Whatever their differences, the partisans of the other four metaethical views all agree that the only facts that exist are facts about our psychology and about the natural world.

T. M. Scanlon seems to provide an alternative to the intuitionist's

non-natural facts. He maintains that we can speak of evaluative judgments being true or false without positing a realm of evaluative facts independent of us. All that is necessary is that there should be standards for correct ways of arriving at evaluative judgments. Evaluative judgments that are arrived at in ways that accord with the standards are true; those arrived at in ways that conflict with the standards are false. Of course, an evaluative claim is not made true by its conformity to any arbitrary standard; it must correspond to *correct* standards of evaluative judgment. As Scanlon recognizes, the judgment that a standard for forming evaluative judgments is correct is itself an evaluative judgment of the kind that the standard applies to.³⁷

It seems that some sort of circularity is involved in Scanlon's account, and indeed I think this prevents his view from constituting a real alternative to intuitionism. Consider some evaluative judgment:

A. Compassion is good.

Suppose that what makes (A) true is the fact that judgment (A) can be arrived at in a way that conforms to the correct standards for arriving at evaluative judgments. I shall abbreviate this fact as follows:

B. Judgment (A) conforms to the right standards.

Now, the judgment expressed by (B) is also evaluative. What makes *it* true? To be consistent, we should say:

C. Judgment (B) conforms to the right standards.

This is also evaluative, and we can ask the same question about it, and so on *ad infinitum*. Scanlon's view thus entails that the truth conditions for any evaluative statement are infinitely complex.³⁸ This strikes me as unbelievable. Furthermore, at no point in the infinite series of conditions is the evaluative element removed; thus, the account does not show that one can accommodate evaluative truths by appealing solely to natural facts.

7.6 The authority of morality

7.6.1 The authority of morality: a rationalist view

Some of our reasons for action derive from desires, specifically, from appetites and emotions. Others—the prudential and impartial reasons—derive from cognitions of facts: the fact that an action best

serves my interests in the long run constitutes a prudential reason for me to take that course of action; the fact that an action would promote the good constitutes an impartial, moral reason for me to take that course of action.

What should one do when reasons of different kinds conflict? This seems like a reasonable and important question, but what does it mean? If 'What should one do?' means 'What is one morally required to do?', then the answer is trivial: obviously, one is morally required to act on one's moral reasons, if any. When we ask the question, intending it to be non-trivial, I think our conception is this: there are different sorts of reasons we have for action, including emotional reasons, prudential reasons, and moral reasons. These are not 'reasons' in different *senses* of the word, but merely different reasons we might have, in some generic sense of 'reason'. These reasons thus can be weighed against each other. So the question arises, when we weigh up all the various reasons we have, what do we have *the most* reason, overall, to do? For instance, do we always have the most reason to be moral, or is it sometimes more reasonable to be immoral? My own view is that prudence trumps desires, and morality trumps prudence. Thus, it is always most rational to act morally. This feature of morality is what I call 'the authority of morality'. The term suggests an analogy to legal authority: for instance, in the American political system, the Constitution takes precedence over statutes, which in turn take precedence over regulations; the Constitution has ultimate authority. This is a normative, not a descriptive claim: it is not that, as a matter of fact, no statutes will ever be enforced that contravene the Constitution, but that such statutes are illegitimate; it is legally required that such statutes not be made or enforced. Analogously, morality takes precedence over prudence, which takes precedence over desires. This does not mean that, as a matter of fact, no immoral actions will ever be intentionally undertaken. It just means that such actions are irrational; one always has sufficient reason against performing such actions.

Why does prudence take precedence over desires? Because a prudential judgment, if formed correctly, *already takes into account* one's present desires. This judgment is not itself a desire, but it is made in the light of a consideration of what one presently desires and how strong those desires are. If one's prudential judgment is that one should act in such a way as *not* to satisfy some present desire, then one has already evaluated some other end as more important than the satisfaction of that desire. To then proceed to weigh the present desire against that judgment would be an unwarranted double-counting of the desire. Notice the asymmetry here: prudential

judgments involve evaluations of desires; desires do not involve evaluations of prudential judgments.

I have not here given a general justification of the demands of prudence—I have not tried to explain why a rational person must assign some weight to his future well-being. I simply assume that. Rather, I have explained why, given one's capacity for prudential judgments, one should not conceive of prudential considerations as being *weighed against* emotions or appetites.

Now why does morality take precedence over self-interest? The answer, I suggest, is structurally the same: because moral judgment, if formed correctly, already takes self-interest into account. This depends on a substantive ethical thesis: that morality is not insensitive to agents' own interests. We are morally entitled to pursue our own interests, within certain limits. For instance, if you pass an accident victim in need of medical attention, you are obligated to call an ambulance, given that the cost to yourself is trivial. But if you witness a mugging, you are not obligated to help the victim by fighting the mugger off, since doing so would put yourself in grave danger.

Thus, if one judges that one is morally required to act against one's own interests on some occasion, that means that one has already weighed one's particular interest in the circumstances against whatever other considerations are at stake, and found one's own interest to be outweighed by those other considerations. To then proceed to weigh one's self-interest against that moral judgment would be an unwarranted double-counting of one's interest. Again, notice the asymmetry: moral judgments involve evaluations of desires and self-interested concerns; desires and prudential judgments do not involve evaluations of moral concerns.

I have not here given a general justification for morality, nor have I refuted ethical egoism. I have merely explained why, given that one has the capacity for moral judgment, one should not conceive of one's moral judgments as being *weighed against* one's selfish motives.

The point here is not that moral considerations are a larger category of which prudential considerations are a subset, nor that prudential considerations are a larger category of which desires are a subset. Desires are not prudential beliefs; they are different in kind. The point is that prudential beliefs include intellectual judgments of the appropriateness and importance of desires. Similarly, prudential beliefs are not moral beliefs.³⁹ The point is just that moral beliefs include intellectual judgments of the appropriateness and importance of self-interested motives. This is what, I claim, gives prudence authority over desire, and gives morality authority over prudence.

7.6.2 The arbitrariness of morality on the Humean view

A Humean can give no similar account of the authority of prudence or morality. We have already seen, in the previous section, the inadequacy of the Humean account of prudence: the Humean must portray prudential reasoning as being, at most, reasoning about how to satisfy a particular present desire, the desire that one's future desires be satisfied. But there is no reason, within the Humean conception, why one must have such a desire, and even if one has it, there is no reason why it should be given any special weight relative to other present desires. The Humean thus has no account of the intuitive notion of the rational authority of prudence.

Likewise, the Humean has no account of the authority of morality, nor any plausible account of why we should attach to morality anything like the importance that we do. A Humean must portray morality as stemming ultimately from some species of desire, such as a desire for the well-being of others. Hume himself posited 'natural sympathy' for other human beings as the source of morality.⁴⁰ All of moral theorizing and moral reasoning is therefore held to be concerned merely with how to satisfy the desires stemming from this particular psychological tendency. But sympathy is just one source of emotion among others; there is no reason for attaching any special weight to the desires it occasions, as opposed to those caused by other emotions. Thus, it would make just as much sense for us to devote a special field of enquiry, complete with complex systems of principles and so on, to figuring out how to satisfy the desires caused by *envy*, *lust*, or *hate*, as it does to devote such a field to sympathy. If my hostile emotions should conflict with my natural sympathy, there is no Humean reason why the latter should generally take precedence. Moral judgments appear to correspond to nothing but an arbitrarily selected subset of our emotions and desires.

How might a Humean try to avoid this consequence? One way might be to portray morality as resulting, not from some particular emotion or subset of emotions, but from a combination of *all* our desires. But if this were so, then we should consider all emotions to give rise to moral reasons, and we should assign moral weight to all desires in proportion to their strength. This is not the case. We do not consider love and hate as providing equally good moral reasons. We do not, for example, think that a sufficiently strong hatred provides a justification for murder; in most circumstances we would take that emotion to provide no reason for action, or even to render actions prompted by it *worse*.

Another approach for a Humean might be to portray morality as

stemming from second-order desires, such as the desire to acquire or lose some other desire. This might make sense of the notion that moral considerations include evaluations of other motives. But this approach faces the same sort of problem as the previous one. Suppose we meet someone who not only hates women, but wants to continue to hate them. We would not characterize his second-order desire as a moral motive, any more than the first-order desire is. Furthermore, for a Humean, one's behavior is always determined by what one *presently wants*, not what one wants to want, what one will want in the future, or anything else. Therefore, relegating moral judgments to judgments about what satisfies second-order desires would prevent them from directly influencing actions. That would be a mistake; moral judgments evaluate actions and outcomes in the world, and it is thus that they can directly conflict with, and perhaps take precedence over, desires and prudential judgments.

7.6.3 The problem of weakness of will

Weakness of the will—the phenomenon in which an individual succumbs to a temptation to act against his own better judgment—is a common experience. It comes in two varieties. In the first, a person acts on some desire, against his own acknowledged best interests. An example is the student who decides to spend the night partying instead of studying, despite her belief that studying would best promote her long-term interests. In the second variety, a person acts on some desire or prudential belief, against his own moral values. An example would be an official who accepts a bribe, despite his awareness that doing so is wrong.

When told of the phenomenon, some people insist that it is impossible for anyone not to do what he thinks best.⁴¹ I assume this view is wrong. Weakness of will is a commonly observed phenomenon. It is the business of theory to accommodate observed facts, not vice versa.

The salient features of the phenomenon include the following:

- i) The agent judges some action, *A*, to be the best.
- ii) The agent feels 'tempted' to do some other action, *B*.
- iii) We say (and, when the agent is us, feel) that the agent's will is 'weak' in some sense; that if the agent were stronger, he would do *A*. Relatedly, it is 'difficult' to do *A* but 'easy' to do *B*.
- iv) The agent typically feels guilt or remorse upon doing *B*.
- v) We judge that the agent is to some degree irrational in doing *B*.
- vi) Although the agent may say such things as 'I couldn't resist', the agent acts voluntarily. He does not, for instance, suddenly lose

control of his muscles and wind up doing *B* in the manner of a muscle spasm. He does not try to do *A* but 'miss' and do *B* by accident. His psyche is not taken over by an alternate personality. No outside force interferes with his decision-making process. Apart from the feeling of conflict between his different motives, the agent's action comes about in exactly the way that intentional actions normally do. He may even engage in conscious planning in order to carry out act *B*, as in the case of the cigarette smoker who, though ostensibly 'trying to quit', deliberately goes to the store to buy a pack of cigarettes.

I have stated the last point with some elaboration because victims of weakness of will have a tendency to make excuses, claiming that they 'can't help' their behavior, that they are 'addicted' (to cigarettes, fatty foods, etc.), and so on. A philosophical grasp of the phenomenon requires resisting that tendency. There is no reason, apart from our desires to avoid incurring blame and to avoid hurting others' feelings by blaming them, to deny that victims of weakness of will can control their behavior. Other than their paradigmatic blameworthiness, weak-willed actions exhibit every feature typical of voluntary, intentional actions.

Now notice the asymmetries implicit in the above features: one is tempted to do *B*; one is not *tempted* to do *A*. Doing *A* is hard; doing *B* is easy. *A* accords with one's 'better judgment'; *B* does not. All of this implies that the reasons for doing *A* are superior to those for doing *B*, but the agent most *wants* to do *B*. A Humean can make no sense of this. For the Humean, when two motives conflict, the only question is, which desire is stronger? If *A* satisfies a stronger desire than *B* does, then the agent will do *A*. If, on the other hand, *B* answers to the stronger desire, then the agent should, rationally, choose to do *B*. In what sense would the agent then be 'weak'? In what sense would he be 'giving in' to a temptation? Why would the agent feel guilt or remorse? And what sense can be made of the idea of one alternative's being 'harder' to choose than the other? Of course, on a Humean view, one might say it is 'difficult' to perform the action for which one has the weaker motives—more precisely, it is impossible.⁴² But this is not the relevant sense of 'difficult', since strong-willed actions are not impossible.

Our ordinary conception of the phenomenon implies that, in cases of weakness of the will, the conflicting motives differ *in kind*. Suppose, for instance, that I am trying to decide between ordering the salmon and ordering the fettuccine for dinner. Suppose the only consideration I have in mind is that of taste. And suppose I deter-

mine that the salmon will taste better. There is no scope for 'weakness of will' here. If I order the salmon, that will not be 'weak'. If I somehow order the fettuccine, that would merely be some bizarre psychological malfunction, not an exercise of strength of will. For a Humean, all motivation is fundamentally of a piece—it is all a matter of present desires of various strengths. Therefore, all cases of conflicting motives are relevantly like the salmon-fettuccine conflict.

The rationalist conception, however, introduces fundamentally different kinds of motivation: our *desires* may conflict with our *judgments* about what to do. That is the essence of weakness of will. And the doctrines of the authority of prudence and of morality explain the asymmetries pertaining to 'weakness', guilt, and irrationality.

The 'problem of weakness of will' is generally taken to be that of explaining how weakness of will is possible. Before answering that, let us first ask: why *wouldn't* it be possible? I contend that the Humean Conception of practical reasons implies that weakness of will is impossible, whereas the Rationalist Conception has no such implication. In a case of weakness of will, the agent acts of his own free will, but he knows that he has sufficient normative reasons for doing something else instead. There are two features of the Humean Conception that rule out such cases. First, as we saw in section 7.5, the Humean account implies that the factors that determine the *normative strength* of reasons are the same as those that determine the *motivating force* of reasons. It would thus be impossible for what an agent is motivated to do to come apart from what it is rational for him to do, as it does in cases of weakness of will.

Second, on the Humean account, all motivations can be compared along a single dimension, where that dimension is the only thing relevant to their capacity to induce action. On the Humean view, it is the strength of the relevant desire that determines its capacity to induce action, and all desires can be compared in terms of their strength.⁴³ The Humean account of motivation therefore strongly suggests a doctrine of *determinism*, that is, that in any given situation (taking into account the agent's motivations), there is only one thing a given agent can do. And this seems, at least on its face, to conflict with the observation that weak-willed actions are freely chosen.⁴⁴ More specifically, it seems that when one describes a scenario in which a person experiences conflicting motivations, and is tempted to act against his better judgment, *either decision is possible*; whichever course of action the agent chooses, it seems that there could have been a situation in which a person had the same reasons available but chose the other course of action.

A Humean might try to explain this by proposing that in such cases, the agent has exactly *equally strong* motivations for each course of action. But if strength of desire varies continuously, then the probability of two distinct desires happening to have exactly the same strength is zero;⁴⁵ therefore, it would not seem reasonable to believe that this is the explanation of weakness of will. And again, the 'exact tie' theory would not explain the various asymmetries implied in the notion of weakness of will, such as the fact that one course of action and not the other is said to accord with one's better judgment, to be rational, to be difficult, and so on.

The Rationalist Conception of reasons does not face the problems that the Humean Conception faces, since the rationalist conception rejects both of the problematic tenets of the Humean view. First, it rejects the idea that the normative strength of reasons is determined by the same factors that determine motivating force. The normative strength of a reason for action is determined, rather, by the objective values—by, for example, how much goodness a given action would produce (or how much one believes it would produce). There is no suggestion on this theory that there must be some psychological dimension of motivating force that this always corresponds to. Second, it rejects the idea that all motivations can be compared along a single dimension of 'strength'. Rather, while *desires* can be compared in terms of strength, our normative *judgments* are a qualitatively different kind of motivation. Judgments have no psychological dimension that can be weighed against the strength of a desire. There is, of course, such a thing as the 'strength' of a belief (that is, the believer's level of certainty), and there is the degree of objective value that one believes some goal to have—but neither of these things is commensurable with the strength of desires. The rationalist conception, consequently, fails to support a deterministic picture of human decision-making. The psychological incommensurability of desires and normative beliefs may explain why, in cases of weakness of will, we think that either decision is possible.

Now I want to clarify two things about this explanation. First, I do not claim that the rationalist conception of practical reasons entails the existence of weakness of will. It is theoretically possible to combine the rationalist conception of reasons with a view in which, in cases of conflicts between different kinds of motives, people always do what they perceive to be morally required—or one in which they always do what they expect to best satisfy their desires. The rationalist conception does not *entail* either indeterminism or the existence of free will, nor does it, by itself, explain 'how free will works'. I have no account to offer of how free will works. What I am arguing is that

the rationalist conception is superior to the Humean conception in that the rationalist conception does not imply that weakness of will is *impossible*, as the Humean conception does.

Second, it is important to understand why, in allowing for cases of weakness of will in the way I have suggested, I am not going back on the thesis of section 7.2. I still maintain that an agent's recognition of normative reasons necessarily produces at least some motivation to act in accordance with those reasons. Recall that there were two central arguments in section 7.2. The first turned on the thought that it must always be *possible* for an agent to do what he rationally ought to do. The second turned on the idea that we are justified *a priori* in interpreting other people, if possible, as being by and large rational. These arguments do not demand a deterministic account of human choice, nor do they imply that the strength of a normative reason must be matched by the strength of the agent's motive in any sense.

7.6.4 How anti-realism undermines morality

We are now in a position to see how belief in anti-realism undermines morality. It undermines both our moral beliefs and our motivation for behaving morally. This fact does not show that anti-realism is false. But it shows that, if anti-realism is false, it is also pernicious.

Some of my anti-realist colleagues seem to find this thesis surprising and hard to accept. Before addressing what they say, it may put things in perspective to observe that it is not only I and other realists who believe this; even some prominent anti-realists have in effect admitted that their work undermines morality, though they don't put it quite that way. Just before attacking the notion of objective value, J. L. Mackie concedes that 'ordinary moral judgments include a claim to objectivity, an assumption that there are objective values in just the sense in which I am concerned to deny this'.⁴⁶ It follows that his own book serves to undermine ordinary moral beliefs. Mackie may believe he can put something better in their place,⁴⁷ but this is quite dubious for the reasons discussed in this section. But it is when anti-realists come to explaining why other people believe in objective morality that they most clearly reveal how their own theories, if accepted, would undermine moral motivation. Consider these explanations provided, respectively, by the anti-realist philosophers J. L. Mackie and Michael Ruse:

We need morality to regulate interpersonal relations, to control some of the ways in which people behave towards one another,

often in opposition to contrary inclinations. We therefore want our moral judgments to be authoritative for other agents as well as for ourselves: objective validity would give them the authority required.⁴⁸

The point about morality . . . is that it is an adaptation to get us to go beyond regular wishes, desires and fears, and to interact socially with people. How does it get us to do this? By filling us full of thoughts about obligations and duties, and so forth. . . . [W]e are then moved to action, precisely because we think morality is something laid upon us. . . . If morality did not have this air of externality or objectivity, it would not be morality and (from a biological perspective) would fail to do what it is intended to do.⁴⁹

It follows that anti-realists, in attacking its objectivity, are working to undermine these important functions of morality.

The point is also clear, I take it, from what I have said above. The only reason anti-realists can offer as to why we should behave morally is that doing so will allegedly help satisfy some pre-existing desire of ours, such as a desire to help other people. If we find that we do not have this desire, or if we find that some other desire happens to be stronger at the moment, then, on the Humean account, it would be irrational to be moral. In any case, the appeal to *moral requirements* would have no special force, relative to any other appeal to desire. Furthermore, if I am right in some of my central claims in this book, both about the meanings of moral claims and about how they motivate us to act, then anti-realism undermines moral motivation in a very direct way: it attacks the beliefs that constitute our reasons for acting morally.

Many anti-realists would resist this conclusion. They would say that moral motivation comes entirely from emotions, so what is needed to bolster morality is simply to train people (applying non-cognitive influences) to have stronger sympathetic, social feelings; and an anti-realist can do this just as well as a realist can. I say that emotions are not enough. People often just do not have the emotions the Humean moralist wants them to have, and they may have no Humean reason (that is, no reason provided by their current desires) for doing anything to change that situation—they may, for example, have no Humean reason for submitting themselves to moral instruction or exhortation.

Consider two illustrations. First, suppose I have just read Peter Singer's argument to show that we have a moral obligation to help

the poor.⁵⁰ According to Singer, each of us is obligated to donate great portions of our income to charities fighting world poverty. Suppose that I find Singer's argument convincing, considered from within the moral standpoint. But then I reflect that if I were to actually do what is 'morally required' according to the argument, I would miss out on a lot of things I presently enjoy—such as going out to eat, going to the movies, and other luxuries. I really like those things. So I'd rather not have to give away my money. So I might think to ask myself: sure, donating to charity is 'morally required', but do I really have to be moral? I turn to the field of metaethics, where I find most experts proclaiming that there are no objective moral requirements. I reflect that this 'moral requirement' Singer proposes is just a matter of social conventions (as the cultural relativists say), or is just something that I have a particular sort of emotion about (as the non-cognitivists and individual subjectivists say), or is just a delusion (as the nihilists say). Believing one of these things, why would I obey the moral requirement? Why would I get out my checkbook and start donating that money, knowing that my own interests will be severely set back in doing so? The only reason Singer can offer me is that by acting morally, I will help satisfy some desire of mine, and perhaps make myself happier.⁵¹ This is pathetically inadequate in light of the kind of sacrifice Singer is calling for. Sure, I have some feelings about the Third World poor—I'm sorry that they're starving, and I feel some satisfaction when I send money to help them. But by no stretch of the imagination are these feelings anywhere close to being strong enough to get me to give away, say, \$30,000 a year. Introspectively, the amount of enjoyment that I derive from sending money to Oxfam is not obviously greater—indeed, it seems much *less*—than the enjoyment I derive, say, from reading a good novel.⁵² And the novel only costs about \$8. So why would I give away thousands of dollars to charity? Granted, my desires are not merely for enjoyment, but the point remains: my desire to help the poor is just not as strong as many of the desires that it competes with. Once I 'realize' that the only reason for obeying moral requirements would be to help me satisfy my personal preferences, I will surely feel entitled to make the judgment that, overall, I'd prefer not to be moral—and to act accordingly. The moral philosopher's telling me that I just don't realize the strength of my sympathetic feelings for other people, and so on, will strike me as about as credible as someone who doesn't know me telling me that I would rather have been a fireman than a professor.

Second, suppose Jon, a married man, is considering whether to

cheat on his wife. He comes to his best friend, Sam, for advice, and the following dialogue ensues:

Jon: [After explaining the situation:] So, do you think I should cheat on Kate?

Sam: Kate would be crushed if she knew you were even considering this.

Jon: Don't worry, she won't find out. [Jon discusses how he's going to make sure Kate doesn't know.]

Sam: Don't you think it's wrong to cheat on your wife?

Jon: Oh, that. That's just an arbitrary convention people made up. I know that a lot of people *don't like* marital infidelity, but why should I listen to them?

Sam: Don't you feel just a little bit guilty for thinking about betraying Kate, after all you've been through together?

Jon: Well, I was feeling guilty for a while. But then I realized that it was stupid to feel guilty. I mean, who says people have to be faithful to their wives? That's just some arbitrary rule someone made up. If I don't want to be faithful, why should I? I admit that I may experience some irrational guilt anyway, but I think it'll still be worth it for all the fun I'm going to have.

Jon has learned the lessons from his philosophy classes well. He may not have followed all the subtleties and complications of the various versions of anti-realism, but he got the gist: we *make up* moral requirements, prompted by some arbitrarily selected subset of our emotions and desires. Our moral sense 'has no authority over those who dissent from its recommendations or even over us if we are inclined to change our minds'.⁵³

I don't think my two examples are unusual. Nor do I see any fundamental flaw in either Jon's or my reasoning in the above examples, apart from the mistake of failing to recognize objective values.

We often hear complaints about the decay of morality in modern society. I do not know whether these complaints are just; I have no statistics on the frequency of immoral acts in different times. But the widespread acceptance of anti-realism in our culture cannot be good for morality. There must surely be people who disregard many of the demands of morality, once they decide that morality is unreal. Of course, ethical intuitionism cannot guarantee moral behavior, any more than it can guarantee correct moral beliefs. No mere philosophical theory can stop people from indulging in weakness of the will.

But at least the theory does not positively encourage such lapses by clothing them in the mantle of practical rationality.

7.7 Why be moral?

Why should we behave morally? The ethical intuitionist's answer is as follows. Our intuitions render us *prima facie* justified in adopting certain moral beliefs. Provided these moral beliefs survive any challenges from potential defeaters, they are justified *tout court*. These beliefs in themselves constitute reasons for action, in virtue of their evaluative content. There is no further question to answer as to why one should be motivated by correct evaluations; it is simply the nature of *evaluation* that it answers questions about what to do. To accept an evaluative judgment *is* to judge something as a reason for action—to accept the judgment that stealing is wrong, for example, is to accept another person's ownership of *x* as a normative reason not to take *x* without his permission.⁵⁴

Some philosophers, with whom I have much sympathy, hold that 'Why should I be moral?' is a nonsense question.⁵⁵ Consider three interpretations of the question:

- i) 'Why should I do what I am morally required to do?' This seems to mean, 'Why should I do what I should do?', which is a nonsense question.
- ii) 'Why should I do the things that are generally considered, and that perhaps I myself have hitherto considered, moral?' Here, the answer would seem to be, 'Because those moral beliefs are by and large true'.⁵⁶ If it is asked why one should accept those beliefs, the answer is that we are justified in believing the contents our ethical intuitions, in the absence of defeaters, as discussed in chapter 5.
- iii) 'Why is it *rational* to be moral?' This interpretation offers the best hope of making sense of the question. If this is what it means, the answer is that moral judgments are themselves products of reason; they are intellectual evaluations of actions, states of affairs, and so on. Furthermore, they are evaluations of a kind that take into account all the relevant reasons, including prudential reasons and one's current desires, as discussed in section 7.6.1. Thus, what we judge moral is what reason evaluates as choiceworthy.

Some readers may be unsatisfied by my answers, for either of two reasons. On the one hand, some may think that, despite my answer to question (iii), there remains a question as to why one should be rational. Others may think that, despite what I have said about it,

question (i) *seems* like a real question. Of course, it sounds better if one replaces 'what I should do' with 'what is morally right' and if one deletes 'should I', thus: 'Why do what is morally right?' I agree that on its face, this *sounds* like a meaningful question. I now want to give a reason for thinking that this appearance is illusory, and that neither 'Why be rational?' nor 'Why do what is morally right?' is a legitimate question.

In order for these questions to be 'meaningful' and 'serious' in the intended sense, there must first of all be something that, in the present context, would count as a satisfactory answer to a question of the form, 'Why do x ?' This answer would presumably say something which is true of x but is not true of other actions that should not be performed. Let that 'something' be F . Suppose, in other words, that 'Because x is F ' is a satisfactory answer to the general philosophical question, 'Why do x ?' I am not assuming that 'Why be moral?' or 'Why be rational?' has a satisfactory answer; I do not assume that moral actions or rational actions have F . I assume only that there is *something* that could be said about these actions that, *if true*, would count as answering those questions. This seems like a reasonable requirement to impose on questions that we are to take seriously.

Yet even if there were such a property F , some philosophers might still say, 'Yes, but why do the F things?' Their question would certainly be grammatically well-formed; moreover, it would probably *sound*, to many at least, as if it were a meaningful question. But by the nature of the case, it could not be a legitimate question; by hypothesis, F is the property that *answers* 'why do so-and-so' questions. Once it has been said that an action is F , therefore, there cannot be any further question as to why one should do it. Conversely, if some such further question needs to be answered, then 'Because x is F ' would not have been a sufficient answer to the question of why one should do x .

We may predict, therefore, that there should be some question or questions of the form, 'Why do the things that have feature F ?' that superficially sound legitimate but are not. These would be questions where the feature cited *already* inherently constitutes (or entails the existence of) a reason for action. What would be the most plausible candidates for such questions? It seems that the most obvious candidates for F are 'moral' and 'rational' (more than one feature may play the role of F)—if *anything* would count as inherently entailing the existence of an adequate reason to do x , it would be that x is right, or that x is rational.⁵⁷

The only other obvious candidate for F is the property of satisfying one's desires. But might I not ask, 'Why should I satisfy my

desires?’ To this, a Humean would respond that the question is illegitimate, because my desiring something already constitutes a reason for me to pursue it. But isn’t it just as clear that something’s being good already constitutes a reason for me to pursue it? If anything, this seems *more* clear. For suppose I have some desires that I myself judge to be desires for the bad. For instance, I feel a sudden, inexplicable urge to jump off a building. Is it really obvious that I then have a reason to jump off the building—that is, something that would make my jumping off the building *prima facie* rational?⁵⁸ Is it obvious, on the other hand, that I have a reason *not* to jump off the building? Of course, I might have both a reason for and a reason against. The point is that it is more obvious that I have a reason not to jump in virtue of knowing that jumping off the building would be bad, than it is that I have a reason to jump in virtue of feeling a desire to jump.

I conclude that such questions as ‘Why be moral?’ and ‘Why be rational?’ are among the clearest candidates for illegitimate ‘why’ questions. They are like the pseudo-question, ‘Why is 10 pounds heavier than 5 pounds?’