

2 Non-Cognitivism

Part 2: Contemporary Responses and Introspective Case

*2.4 Hare on moral truth

Most of the arguments discussed in the preceding section have been known in some form to philosophers for over 40 years, yet there remain non-cognitivist philosophers. In this and the following two sections, I consider how some prominent non-cognitivists respond to some of these problems.

To begin with, R. M. Hare has addressed the problem of moral truth, endeavoring to show that his version of prescriptivism can preserve a notion of truth and falsity for moral statements. How? Two distinct proposals can be found in his work. First, Hare proposes that moral statements have a descriptive, proposition-expressing meaning in addition to their prescriptive meaning. It is this descriptive meaning that enables them to be true or false. Hare seems to think of the descriptive meaning of a moral statement as equivalent to that of a statement of the non-moral facts that are the speaker's reasons for making the moral judgment.¹⁶ For example, suppose Jon has a comprehensive moral theory, according to which the morally right action in any situation is always the action that produces the greatest amount of happiness in the world. Jon says, 'Eating meat is wrong'. According to Hare, Jon's statement has a descriptive meaning and a prescriptive meaning. The descriptive meaning is, 'Meat-eating decreases the amount of happiness in the world'. The prescriptive meaning is something like, 'Don't eat meat'.¹⁷ The original statement, 'Eating meat is wrong', can be assessed as true or false by reference to whether eating meat in fact decreases the amount of happiness in the world.

I think this is a false conception of meaning. When one makes a moral claim, the meaning of the claim is one thing; one's reason for it is another, and the latter is generally not part of the former. For example, suppose that George favors invading Iraq. His reason is that he thinks Iraq has weapons of mass destruction, and he thinks this an adequate justification for invasion.¹⁸ A second individual, Tony, also thinks we should invade Iraq. But Tony does not believe Iraq has weapons of mass destruction. Tony's reason for favoring invasion is that Saddam Hussein is a cruel dictator who has caused much suffering for the people of Iraq. George and Tony have the following exchange:

George: We should invade Iraq.

Tony: I agree. We should invade Iraq.

On Hare's view, Tony is confused in thinking that he agrees with what George said. For though Tony may be said to agree with George's plan of action, part of the meaning of George's statement was that Iraq has weapons of mass destruction, which Tony does not believe. So Tony actually doesn't agree with the statement as a whole. But this is obviously false. While George may have been thinking that Iraq has weapons of mass destruction, and while this may be part of the reason why he said what he did, George did not, in the above exchange, actually say that Iraq has weapons of mass destruction, nor did he say anything which implies this.

I do not know whether Hare would actually consider 'Iraq has weapons of mass destruction' to be part of the meaning of George's statement. Perhaps Hare would instead require us to formulate George's overall, comprehensive ethical theory. Unfortunately, most people have no such theory. But suppose we are lucky in this case and George has a comprehensive ethical theory. Let us say his theory is utilitarianism. Then, Hare might say, the descriptive meaning of George's statement is, 'Invading Iraq would maximize total happiness in the world'. While this might seem less absurd than the previous suggestion, it is still false: in saying merely that we should invade Iraq, George does not say that doing so would maximize happiness, nor does he say anything which implies this; he simply has not said anything, so far, about why he thinks we should invade Iraq. And if Tony should have a competing ethical theory, this would not preclude him from agreeing with George's conclusion by saying, 'Yes, we should invade Iraq'.

Consider a second example, even more damaging for Hare's conception of moral truth. As in the first case, George thinks we should invade Iraq because Iraq has weapons of mass destruction. Jean-Pierre also believes that Iraq has weapons of mass destruction and agrees with George on all other non-moral matters. But Jean-Pierre and George have different fundamental moral values, as a result of which Jean-Pierre does not take these facts to justify invasion. George says, 'We should invade Iraq'. According to Hare, Jean-Pierre should judge that statement to be true.^[19] For it is the descriptive meaning of the statement that determines truth or falsity, and Jean-Pierre agrees with George on descriptive matters.

Thus far, I have been talking as though Hare would rely on the speaker's moral standards to determine what the descriptive meaning of a moral statement is. In fact, however, Hare adverts to the 'commonly accepted' standards 'within a given society'.^[20] Thus, if George and Tony belong to the same society, they would wind up (thankfully) meaning the same thing by 'We should invade Iraq'. But it is easy enough to revive the above problems by simply supposing George to belong to a different society from Tony or Jean-Pierre. And the reliance on socially accepted standards just creates a further problem for Hare's account, that of how to deal with people who disagree with their society's moral standards. Such people would, on Hare's view, be forced to call moral judgments made on the basis of the prevailing moral standards 'true'. And when such moral dissidents called something 'good', part of what they would be saying would be that the thing satisfied society's moral standards. It seems clear that, if anyone's moral standards are involved in determining the descriptive meaning of a moral claim, it ought to be the speaker's standards.

I take it that Hare's first effort to secure truth and falsity for moral statements fails. It is, in fact, the product of an unfortunately common fallacy to which philosophers have yet to be fully immunized: that of confusing a speaker's substantive background beliefs with the meanings of his words. A special case of this is confusing a person's reasons for saying something with the meaning of the statement. When we express a moral judgment, we are not thereby reporting the factual grounds we have for the judgment.

Hare makes a second attempt to secure moral truth, using a non-cognitivist analysis of 'true'. Hare suggests that in some cases, we might call a moral statement 'true', meaning thereby simply to endorse it: 'it will not be a statement that we endorse it, but an expression of endorsement'.²¹

One problem with this suggestion is that the word 'true' is not used this way in English. In English, 'true' is only applied to (things expressing) propositions. Thus, I cannot say, 'It is true that please pass the salt', even if I endorse the imperative to pass the salt. I cannot say, 'It is true that hurray for the Broncos', even if I endorse the sentiment behind the sentence 'Hurray for the Broncos'. Hare would thus have to claim that evaluative statements work differently from all the paradigmatic non-cognitive utterances in that, for some reason, one can apply 'true' to them.

In addition, Hare's proposed non-cognitivist interpretation of 'true' faces many of the same problems as a non-cognitivist analysis of 'good'. The sentence 'x is true' bears all the linguistic marks of a factual assertion. We can speak of 'the proposition that x is true' and 'the truth of my claim'; we can say, 'I believe x is true', 'I hope x is true', 'I wonder whether x is true', 'If x is true, then so is y', 'Do you swear to tell the truth?' and 'There is many a true word spoken in jest'.²² All of these are hard to make sense of on a non-cognitivist construal of 'true'. And the considerations usually adduced in support of non-cognitivism about 'good'—such as the alleged unworkability of alternative theories and the action-guiding character of moral judgments—will not aid the non-cognitivist about 'truth'.

A better proposal, perhaps closer to what Hare intended, is the redundancy theory of 'truth': 'it is true that' is simply redundant. 'It is true that pleasure is good' just means 'Pleasure is good', just as 'It is true that cows have four stomachs' just means 'Cows have four stomachs'. This gives us a principled reason for holding 'It is true that pleasure is good' to be non-cognitive and 'It is true that cows have four stomachs' to be cognitive, without having to say that 'true' shifts its meaning.

The problem with this suggestion is that it can make no sense of statements in which 'true' or 'truth' is used without specifying what is or might be true, such as 'Do you swear to tell the truth?' and 'Many a true word is spoken in jest'. Neither a redundancy theory nor an endorsement theory of 'true' handles these examples. Nor can we interpret these statements along the lines of 'Do you swear to say what you endorse?' and 'Many a word endorsed by the speaker is spoken in jest', since either of these interpretations would abandon the non-cognitivist treatment of truth for a subjectivist theory, treating truth merely as the property of being endorsed by the speaker.²³ Note that sentences of this kind may just as well advert to moral truths as non-moral ones—'Jon spoke the truth' may be said in the same sense when Jon's salient remark was 'Cats have teeth' as when his salient remark was 'Burning cats is wrong'.

Some of Hare's own statements constitute counter-examples to non-cognitivism both about ethics and about truth. Immediately following Hare's proposal of the endorsement theory of moral 'truth', we find these remarks:

So we cannot make standards acceptable simply by accepting them; rather, in saying that they are acceptable we are expressing our acceptance of them; and whether we should accept them is a matter for rational thought. . . .

We have to do this thinking, and there are ways . . . of doing it well.

If we do it well, we shall attain knowledge of the truth of moral statements. . . .²⁴

'Acceptable', 'should', 'have to', and 'well' are all evaluative expressions. It is unclear what any of the above statements are supposed to mean on a non-cognitivist construal of evaluative expressions. Nor can I understand what 'We shall attain knowledge of the truth of moral statements' is supposed to mean if 'truth' is a non-cognitive term serving only to express endorsement of a given statement. It is tempting, particularly in view of the rest of Hare's views, to interpret this as meaning something like, 'We shall attain knowledge of which moral statements would be endorsed by all rational beings after sufficient reflection'. But to accept this interpretation would be to drop the proposed non-cognitivist theory of truth in favor of a cognitivist one in which 'true', applied to evaluative statements, means 'would be endorsed by all rational beings after sufficient reflection'. But if those are the truth-conditions for evaluative statements, then we might as well just analyze moral terms as referring to what rational beings would endorse, and abandon non-cognitivism for subjectivism. I conclude that Hare has failed to articulate a legitimate sense of moral truth consistent with his metaethics.

*2.5 Gibbard's factual-normative worlds

The latest in anti-realist technology, Allan Gibbard's 'factual-normative worlds' are abstract objects designed to help explain the meanings both of simple evaluative statements and of complex sentences with evaluative components. A factual-normative world is a combination of a complete specification of the way the world is, w , with a complete set of rules for acting and for approving/disapproving of things, n . We can represent a factual-normative world as an ordered pair, $\langle w, n \rangle$.

The meaning of an ordinary factual statement with no evaluative terms can be represented by a set of possible worlds—roughly, the set of all ways the world might be that would make the statement true.[26] Similarly, Gibbard thinks, we can represent the meaning of a statement containing evaluative terms by a set of factual-normative worlds—the set of all factual-normative worlds in which the statement holds. A statement 'holds' in a factual-normative world $\langle w, n \rangle$ if, roughly, someone who thought the world was like w and who endorsed norms n would accept the statement. For example, suppose you have just stolen a piece of candy from an infant, and a friend says: 'If stealing that candy was wrong, God will punish you'. Now imagine two possible worlds, w_1 and w_2 , and two possible complete sets of norms, n_1 and n_2 , where:

In w_1 , God punishes you.

In w_2 , God does not punish you.

n_1 permits the candy theft.

n_2 forbids the candy theft.

From these worlds and these norm-sets, four factual-normative worlds can be formed: $\langle w_1, n_1 \rangle$, $\langle w_1, n_2 \rangle$, $\langle w_2, n_1 \rangle$, and $\langle w_2, n_2 \rangle$. The statement 'If stealing the candy was wrong, then God will punish you' holds in $\langle w_1, n_1 \rangle$, $\langle w_1, n_2 \rangle$, and $\langle w_2, n_1 \rangle$; it does not hold in $\langle w_2, n_2 \rangle$. So $\langle w_2, n_2 \rangle$ is the only one of the four worlds that the speaker's statement rules out. Gibbard's view can be applied equally well to more complex statements and more general statements about wrongful actions (for instance, 'If it is wrong to do something, then it is wrong to get your little brother to do it').

What does the device of factual-normative worlds accomplish? It gives Gibbard a way of saying that the sort of statements that worried us in section 2.3—such as 'If stealing is wrong, God will punish you'—are meaningful. Importantly, Gibbard need not posit any shift in the meaning of 'wrong' from 'Stealing is wrong' (asserted by itself) to 'If stealing is wrong, God will punish you'. Nor must he give separate theories of the meanings of different kinds of sentences containing evaluative terms; there will be general rules relating the meanings of complex sentences to those of their simpler parts, which are the same for evaluative and non-evaluative statements—for example, the non-cognitivist won't have to say that 'if . . . then' has different meanings depending upon whether an evaluative or a descriptive statement is plugged in for the antecedent.

To gain these appealing advantages, we need only draw the link between factual-normative worlds and meanings. What exactly does a set of factual-normative worlds have to do with what I mean when I make some statement containing 'wrong'? Gibbard thinks that to explain the meaning of a statement is to explain the state of mind that it expresses. His view is not, of course, that I actually have the relevant set of factual-normative worlds in mind—possible worlds are typically too complex for any human mind to fully conceive, to say nothing of the infinite sets of factual-normative worlds to which Gibbard's theory adverts. Rather, Gibbard's view is that we have two sorts of mental states, beliefs and attitudes

of norm-acceptance, and mental states of each of these kinds can be said to rule out certain factual-normative worlds. This 'ruling out' is a matter of objective logical relations, so my beliefs and norms might rule out things that I don't realize they rule out (because I have not followed out all the logical consequences of my beliefs and norms). Here are a couple of examples: If I believe that it is going to snow tomorrow, then my belief rules out every factual-normative world in which it doesn't snow tomorrow. If I accept a norm of not killing people, then I rule out any factual-normative worlds in which killings of persons are permitted. Notice that this norm not only rules out some competing norms, but also rules out some fact-norm combinations. For instance, my no-killing-people norm rules out the following combination: the

factual claim that selling cigarettes kills people + the norm of approving of cigarette sales. So any factual-normative world including both of those things would be ruled out; however, a factual-normative world that approves of cigarette sales but contains the factual claim that selling cigarettes does not kill people would not be ruled out.

When you make a statement, you express a state of mind—a set of beliefs and norm-acceptances—that rules out certain factual-normative worlds. The set of factual-normative worlds in which your statement holds is just the set of worlds that your mental state does not rule out. That is what the set of worlds we have been discussing has to do with meaning.

It may seem that all of this solves the Frege-Geach problem. Gibbard seems to think so. But in fact it does not touch the core problem.

Let us begin with a thought experiment in the philosophy of language. It is commonly (though not universally) accepted that there are things called 'propositions' that our beliefs are about. If you believe it will rain tomorrow, then your belief is directed at the proposition that it will rain tomorrow. If you believe that Snoopy will get wet, then your belief is directed at the proposition that Snoopy will get wet. Now imagine that we met a philosopher who held the interesting thesis that there are no compound propositions. For example, there is a proposition that it will rain tomorrow and a proposition that Snoopy will get wet, but there is no such thing as the proposition that if it rains tomorrow, then Snoopy will get wet. This philosopher would run into trouble, because there certainly is such a thing as believing that if it rains tomorrow then Snoopy will get wet. The philosopher would have a hard time accounting for this, since he would have available no proposition that this belief could be directed at. He admits the state of believing that it will rain tomorrow, as well as the state of believing that Snoopy will get wet. So he might try from those two beliefs to somehow construct the compound belief. But the person who believes that if it rains tomorrow Snoopy will get wet, need not believe either that it will rain tomorrow or that Snoopy will get wet. So it appears the philosopher would lack the resources to do what he needs.

Now consider a different (alleged) belief: the belief that if it was wrong to steal the candy, then God will punish you. With respect to this belief, the non-cognitivist is in the same position as the hypothetical philosopher of the last paragraph. The non-cognitivist recognizes a state of accepting an anti-candy-stealing norm, which corresponds to the expression 'it was wrong to steal the candy', and he recognizes a state of believing that God will punish you, which corresponds to the expression 'God will punish you'. But there can be no mental state that corresponds in a similar way to the whole sentence, 'If it was wrong to steal the candy, then God will punish you', because someone who sincerely asserts that sentence need neither accept the anti-candy-stealing norm, nor believe that God will punish you. Gibbard's only psychological resources are beliefs and norms, but there is no possible combination of beliefs and norms that can be taken as corresponding to the conditional.

To put the matter another way: Gibbard wants to explain the meaning of the conditional, 'If it was wrong to steal the candy, then God will punish you', by reference to the set of factual-normative worlds that the speaker's state of mind rules out. As we saw above, the correct set of worlds to rule out is just $\{\langle w_2, n_2 \rangle\}$. But there is no possible combination of mental states—using the resources at Gibbard's disposal—that rules out just this world. The relevant possible mental states, on Gibbard's theory, are:

- Believing that God will punish you.
- Believing that God will not punish you.
- Having no opinion about whether God will punish you.
- Approving of the candy-stealing.
- Disapproving of the candy-stealing.

- Being neutral towards the candy-stealing.

Nine consistent combinations of mental states can be formed from these. None of them has the desired result of ruling out just $\langle w_2, n_2 \rangle$. For example, suppose that I believe God will punish you and I approve of the candy-stealing. This doesn't work, since it rules out $\langle w_2, n_1 \rangle$. In fact, no combination that includes believing that God will punish you works, since any such combination rules out $\langle w_2, n_1 \rangle$. No combination that includes believing that God will not punish you works either, since it will rule out $\langle w_1, n_1 \rangle$. No combination that includes approving of the candy-stealing works, since it will rule out $\langle w_1, n_2 \rangle$. No combination that includes disapproving of the candy-stealing works, since it will rule out $\langle w_2, n_1 \rangle$. This leaves only the combination having no opinion about whether God will punish you and being neutral towards the candy-stealing. But this fails to rule out $\langle w_2, n_2 \rangle$.

The problem here is the separation between belief states and norm-acceptance states, which is made inevitable by their being different kinds of mental states with no common objects. It is not in general possible to construct an attitude towards a compound proposition from attitudes towards the separate components of the proposition.

This is why Gibbard cannot in general supply meanings for compound statements in which one component is evaluative and the other non-evaluative.

Notice that the cognitivist has no problem here. The cognitivist need not try to construct the belief that if stealing the candy was wrong then God will punish you out of separate opinions about the candy theft and about God's possible future punishment. Instead, the cognitivist has a single object, the proposition that if stealing the candy was wrong, then God will punish you, and a single mental state (a belief) that can be directed towards that proposition.

*2.6 Blackburn's solutions

A suggestion due to Simon Blackburn might help Gibbard. However, if it does, it would render the discussion of factual-normative worlds unnecessary; hence, I will discuss this as a separate proposed solution to the Frege-Geach problem. Blackburn's suggestion is that in uttering a conditional statement, one expresses a dispositional mental state: a disposition to accept the consequent, if one were brought to accept the antecedent. So if I say, 'If stealing the candy was wrong, then God will punish you', I am expressing my disposition to believe that God will punish you, if I should come to disapprove of the theft.²⁸ Presumably, Blackburn would say the same about conditionals that don't contain evaluations (the ability to do so would be an important virtue of this theory). So 'If it rains tomorrow, Snoopy will get wet' expresses a disposition to accept that Snoopy will get wet, if one becomes convinced that it will rain tomorrow.

Here is a story to illustrate why that's wrong. A psychiatrist has a patient who thinks he (the patient) is dead. The shrink says to the patient: 'All right. You know that the blood of dead people coagulates, right?' Patient: 'Yes.' 'So if you're dead, then if I prick your finger, the blood will not flow out, right?' 'Yes.' The shrink produces a needle, pricks the patient's finger, and the blood flows out. Shrink: 'What do you say now?' Patient: 'Well, I guess I was wrong. I guess the blood of dead people doesn't coagulate.'

In this example, the psychiatrist attempts to engage the patient with the following inference:

If Patient is dead, his blood has coagulated.

Patient's blood has not coagulated.

Therefore, he is not dead.

The patient initially accepts the first premise. He is brought to accept the second premise. But instead of thereupon accepting the conclusion, he simply revises his belief in the first premise.

Similarly, if a person accepts 'If A then B ', he need not be disposed to accept B upon being convinced of A ; he might instead be disposed, upon becoming convinced of A , to stop believing 'If A then B '. In some cases, this would be perfectly rational. For instance, suppose I believe that since Ann is a reliable moral judge, if Ann thinks a person is evil, then that person is evil. But when Ann unexpectedly declares that Mother Theresa was evil, I do not accept that Mother Theresa was evil; instead, I give up the belief that if Ann says someone is evil, then they are. This shows that we cannot interpret accepting a conditional as being disposed to accept the consequent if one should come to accept the antecedent.

A qualification that appears in Blackburn's statement of the theory might be thought to help here:

Suppose I hold that either John is to blame, or he didn't do the deed. Then I am in a state in which if one side is closed off to me, I am to switch to the other—or withdraw the commitment.³⁰

The italicized phrase seems to let Blackburn escape from my counter-example. Blackburn could say that when one sincerely asserts a conditional, one is disposed, upon accepting the antecedent of the conditional, to either accept its consequent or withdraw one's commitment to the conditional. In the preceding two examples, it is the latter that happens.

Remember that the original point of Blackburn's proposal was to enable the non-cognitivist to explain what it means to accept a conditional statement with an evaluative antecedent. While it may be true that a person who accepts 'If \$p\$ then \$q\$' is disposed, upon being convinced of \$p\$, to either accept \$q\$ or give up 'If \$p\$ then \$q\$', this cannot be taken as an explanation of what 'If \$p\$ then \$q\$' means. The effect of including the phrase 'or withdraw the commitment' within the characterization of what the commitment is supposed to be a commitment to, is to empty the supposed commitment of content. It is as if I said, 'I hereby promise to either pay you \$10,000 or go back on this promise'. A commitment to (either perform act \$A\$ or withdraw this commitment) is equivalent to a 'commitment' to (either perform \$A\$ or not), which is no commitment at all. Similarly, a disposition to (either accept \$q\$ or stop having this disposition) is hard to distinguish from a disposition to (either accept \$q\$ or not).

To put the problem differently, the claim that a person has some disposition ought to imply some predictions about what the person would be likely to do (in non-dispositional terms) in at least some circumstances; otherwise, having the disposition would be indistinguishable from not having it. That is, claims about one's dispositions must entail determinate predictions regarding actions apart from the 'actions' of acquiring and losing dispositions. Blackburn's hypothesized disposition violates this constraint. Once 'or withdraw the commitment' is added to the characterization of what one is disposed to do, there is no longer any prediction as to what (non-dispositional) states a person would be likely to go into in any circumstances. It is now perfectly consistent with Blackburn's analysis that a person having the supposed disposition would behave in exactly the same way (characterized non-dispositionally) in every possible circumstance as a person lacking the disposition. The reason why Blackburn's qualification lets him escape from my counter-example, then, is that it empties his analysis of content.

While the initial theory of the meaning of conditionals may seem plausible at first glance, Blackburn's view of disjunctions, as expressed in the above quotation, does not. He contends that to accept a disjunction is to be disposed to accept either disjunct if the other one should be closed off to you. For example, I accept the statement 'Either there are unicorns living on the planet Jupiter, or cows have four stomachs'. (This is true.) According to Blackburn, this means that I am disposed to accept that there are unicorns on Jupiter if I should find out that cows don't really have four stomachs. But nothing remotely like that is the case.

Blackburn earlier put forward a slightly different theory, before he came to the one we've been discussing.³¹ He proposed that compound statements with evaluative components express a kind of second-order evaluation—roughly, they express approval towards the sort of dispositions we've been discussing. So 'If stealing the candy was wrong, God will punish you' expresses the speaker's approval towards the having of (a disposition to accept that God will punish you, if one comes to disapprove of the candy theft). The difference between Blackburn's earlier and later views is that in the earlier view, he says we express approval towards certain dispositions, while in the later view, he says simply that we express the dispositions themselves. Does this difference affect the force of my objections? No—in essence, the same two objections apply. First, accepting a conditional does not commit one to approving of (a disposition to accept the consequent if one accepts the antecedent). One might instead approve only of a disposition to give up the conditional upon adopting the antecedent. Second, Blackburn's account cannot handle disjunctions. 'Stealing is wrong' logically implies 'Stealing is wrong, or unicorns live on Jupiter'. But surely my disapproval of theft does not commit me to approving of a disposition to accept that unicorns live on Jupiter, if one should come to approve of theft.

Even setting aside these points, Blackburn's proposals thus far could at most provide an account of what goes on in sentences like 'If pleasure is good, then cinnamon ice cream is good' and 'Either Simpson didn't kill Brown, or the verdict was unjust'. We still don't know how, on a non-cognitivist view, to understand such sentences as the following:

1. I believe that cinnamon ice cream is good.
2. I wonder whether abortion is wrong.
3. I hope that I did the right thing.

4. People don't always do what's right.

5. Don't reward people for doing bad things.

Let us discuss each of these five sentences in turn.

(1) 'I believe that cinnamon ice cream is good.' Blackburn offers: '[B]elieving that X is good or right is roughly having an appropriately favorable valuation of X'.³² Notice that this is not what 'believe' normally means. Blackburn is positing an ambiguity in the word 'believe': when I say I believe pleasure is good, according to Blackburn, I actually am not saying that I believe anything in the sense in which I would say I believe it's going to rain tomorrow.

Consider one way of testing for ambiguity; I call it the joke test. Someone says, 'Stephanie has very fair skin. Her skin is even lighter than a feather'. The second sentence sounds funny because it treats two meanings of 'light', light as opposed to dark and light as opposed to heavy, as if they were the same. Now contrast the remark, 'Jerry is very pro-life. He believes that abortion is wrong even more strongly than that 2+2 is 4'. This remark is not funny at all. But if 'believe' were ambiguous, such that one could not 'believe' that abortion is wrong in the same sense that one can 'believe' that 2+2 is 4, then the second example should sound funny in the same way as the first.

(2) 'I wonder whether abortion is wrong.' Blackburn says: '[W]ondering whether X is good or right is wondering what to do/ what to admire or value'. But though I wonder whether abortion is wrong, I thankfully am not wondering whether to have an abortion. Am I wondering whether to condemn abortion? Perhaps, but that would be another matter: wondering whether to condemn abortion would be, if anything, wondering whether condemning abortion is right, not wondering whether abortion is right. Blackburn's proposal becomes still less plausible when applied to wonderings about the

past or about other people: if 'I wonder whether it is right for me to do A' means 'I wonder whether to do A', then 'I wonder whether it was right for Truman to bomb Hiroshima' means . . . what?

(3) 'I hope I did the right thing.' Blackburn offers no help on this one, and I see no interpretation of this sentence paralleling his treatments of the previous cases—nor any other non-cognitivist interpretation, for that matter.

(4) 'People don't always do what's right.' Blackburn does not discuss this one, but taking a cue from his treatments of (1) and (2), we might try this paraphrase: 'People don't always do what is to do'. But this makes no sense. We might say that the sentence expresses disapproval towards some things that people do, but it is hard to see someone expressing disapproval just towards 'some things' but not towards anything in particular.

(5) 'Don't reward people for doing bad things.' We might try interpreting this as expressing disapproval towards a practice of rewarding people for doing things that one disapproves of. But this would be to treat 'bad thing' as a descriptive term equivalent to 'thing that one disapproves of', which is to abandon non-cognitivism in favor of subjectivism. The proposal that the sentence expresses disapproval towards the practice of rewarding people for doing things that the speaker disapproves of, similarly reverts to subjectivism.

In contrast, the cognitivist view is simple and straightforward. It does not require one to posit any ambiguities in words like 'believe', nor to think up eccentric semantic theories for each of the 'problem' sentences. On the cognitivist story, all of these sentences are easily understandable and mean just what they appear to mean.

***2.7 Timmons' assertoric non-descriptivism**

Since its inception, non-cognitivism, in its efforts to evade objections, has developed ever more resemblance to moral realism. Simple non-cognitivists such as Ayer just said moral statements don't assert anything, don't express beliefs, and can be neither true nor false; nor are there any moral properties or facts. Later non-cognitivists, such as Hare and Blackburn, tried to soothe common sense by letting us apply the words 'true' and 'false' to moral statements and speak of valid arguments involving them, while still maintaining the rest of the doctrine—that is, that moral statements do not assert propositions or express beliefs, nor are there any moral properties or facts. This philosophical trick required some reinterpretation of 'true' and 'false', but that was considered acceptable. The object of the game, for many philosophers, is to devise interpretations of common English words and phrases that make most of the statements ordinary people would accept come out as true—or at least as appropriate things to say—while still holding on to one's preferred metaphysical views (in this case, the view that there are no objective values). One of the rules of the game is that you never call your interpretations reinterpretations; you always say you are giving the current actual meanings of the relevant words. Bishop

Berkeley inaugurated this practice, attempting to show that his denial of the existence of matter was perfectly consistent with common sense,³³ and contemporary philosophers have raised it to an art form.

Returning to the history of non-cognitivism, the theory was not yet reckoned successful, because ordinary people will accept such statements as, 'I believe eating cows is wrong', 'Jon asserted that eating cows is wrong', 'The fact that the war is unjust explains why people are protesting it', and 'Slavery has the property of injustice'. Non-cognitivism, even of the Hare-Blackburn variety, still had to regard those statements as false, if not nonsensical.

Mark Timmons has recently produced the next, and presumably last stage in the non-cognitivist assimilation to realist ways of talking. His 'assertoric non-descriptivism' repudiates—at least verbally—nearly all of the theses that traditionally defined non-cognitivism. He allows that moral statements are assertions, express beliefs, and can be true or false. There are moral properties and facts. Timmons even repudiates relativism: the moral truth does not depend on my or anyone else's attitudes. However—this is the part that makes Timmons an anti-realist—moral statements are not descriptive, but evaluative. 'Descriptive' assertions are ones that purport to describe states of affairs in the world, and so require for their truth that the world be a certain way; in contrast, 'evaluative assertions' function to guide action and to express a mental state that consists of 'taking a moral stance on some issue'.³⁴

If this combination of theses strikes you as inconsistent, you may be laboring under misapprehensions concerning the meanings of some of the words in the preceding paragraph, such as 'true', 'assertion', 'belief', 'property', and 'fact'. Begin with 'true'. The traditional view of truth is the correspondence theory: for a statement to be True is for it to correspond to the Facts in a mind-independent, outside World. For convenience, let us adopt the notational convention of using all-caps when referring to the denizens of this mind-independent world; so OBJECTS, PROPERTIES, and FACTS are the sort of things that occupy the WORLD according to the correspondence view, and that our statements need to correspond to.³⁵

Timmons, following some recent philosophical trends, prefers a 'minimalist' interpretation of truth. For a statement to be true, he says, is just for it to be 'correct' to assert it, according to whatever the norms are for the class of statements to which it belongs. For some kinds of statements, the rules we have for when they are correct include that the external WORLD must be a certain way. But the rules for other kinds of statements—including evaluative ones—do not make such demands on the WORLD. When I say, 'It is true that eating meat is wrong', I am really just saying, 'Eating meat is wrong'—the latter statement does not make any demands on the WORLD (again, it just expresses my moral stance), so neither does the former.³⁶

Next, 'assertion': for Timmons, any sentence that has all the linguistic markers of an assertion, of the sort discussed in section 2.3, is ipso facto an assertion; that is all there is to being an assertion. So moral statements are assertive. Likewise, there are moral beliefs, because any mental state that can be expressed by an assertion automatically counts as a 'belief'. There are moral facts, because all that is required for there to be a 'fact' (but not a FACT) is that there be a true assertion. There are moral properties (not PROPERTIES), because all that is required for there to be 'property' is that there should be an adjective that we correctly apply to things. Thus, Timmons gets to say nearly all the things that realists customarily say. His statements only differ from those of the moral realist when using block letters: then he says there are no moral FACTS or PROPERTIES, nor are any moral statements TRUE.

That is just an overview, but it suffices to set the stage for my criticisms. I think that the best way of understanding what is going on in Timmons' work also shows what is unsatisfying about it. In essence, Timmons' view comes to the same as that of the traditional non-cognitivists, but with the following stipulations about word usage: Timmons extends the word 'assertion' to cover both what they called 'assertions' and what they called 'emotive expressions'; he extends the word 'belief' to cover both what they called 'beliefs' and what they called 'moral attitudes'; and he extends the word 'true' to cover both what they called 'true beliefs' and what they called 'admirable attitudes'.³⁷ His view that there are genuine assertions that are not descriptive initially sounds radical, even paradoxical; but properly understood, he is not saying that moral statements are assertions in the sense that earlier theorists understood the word 'assertion'. Of course, he would say that earlier thinkers misunderstood the English word 'assertion', but let's leave that aside. If Assertoric Non-Descriptivism differs only verbally from traditional non-cognitivism, then it seems unlikely that it offers any real advance on the problems faced by non-cognitivism.

Indeed, it does not. Timmons' treatment of the Frege-Geach problem is sufficiently like Blackburn's, and his treatment of moral truth sufficiently like Hare's, that I won't discuss them separately here. Instead, I shall raise two objections specific to his views.

The first concerns moral error. Unless we are dogmatic, we all admit that some of our moral views could be wrong. Such statements as, 'I think cloning is good, but I might be mistaken' are often apposite. The realist understanding of that remark is straightforward, but how can an anti-realist understand it? How can I be mistaken, if there are no moral facts independent of my attitudes? Timmons views this as the most serious problem facing anti-realism. His solution is that an admission of moral fallibility expresses an attitude towards one's own moral attitudes, namely 'an attitude of openness to the possibility that [one's moral outlook] might evolve into a different moral outlook as a result of new and persuasive arguments, as well as new experiences'.³⁸

This clearly is not right. What Timmons describes might be what is expressed by 'I may come to believe my current outlook to be mistaken', but not what is expressed by 'My current outlook may be mistaken'. Consider a slightly different remark: 'My moral views may be mistaken in ways I shall never discover'.³⁹ This, too, makes perfect sense, but it evidently does not express an attitude about the possibility of my changing my moral views in some way in the future. Or suppose I remark, 'I may wrongly come to adopt a different moral outlook as a result of new and persuasive arguments'. This makes sense, and it entails that my moral outlook may evolve in the way Timmons describes, but it does not imply that my current outlook may be mistaken. Thus, one can avow moral fallibility without expressing openness to the possibility of one's outlook evolving, and one can express openness to the possibility of one's outlook evolving without avowing moral fallibility.

Furthermore, compare the two sentences,

I think eating meat is wrong, but I might be mistaken.

I think Napoleon was defeated in 1815, but I might be mistaken.

The latter sentence, presumably, just talks about how the mind-independent world might be, rather than about possible future changes in my beliefs. So Timmons must think that 'might be mistaken' has different meanings in the two sentences. If so, then this should be nonsensical:

So there are two things I might be mistaken about: the moral status of meat-eating, and when Napoleon was defeated.

But that follows from the two preceding statements, and it isn't nonsensical.

My second objection to Timmons' account is more fundamental. Timmons wants to show that common sense is not committed to moral realism, by offering interpretations of how language works that allow us to say the things we normally say without being committed to the existence of evaluative FACTS. The ploy fails because his interpretations are simply not what we mean by 'truth', 'fact', and so on. I propose to show this by arguing that, if they were what we meant, then we could not understand Timmons' exposition of his own theory; in particular, we would see no distinction between assertoric non-descriptivism and moral realism.

Timmons has only the English language to use in explaining his theory. He can introduce new, technical terminology, but then he must explain, in English, what his technical terms mean. Now, the distinction he posits between properties and PROPERTIES is crucial to explaining his theory and how it differs from moral realism; roughly, Timmons isn't a realist because he believes in moral properties but not moral PROPERTIES. How does he explain this distinction? Here is the entire passage introducing it:

According to the correspondence theory, truth is a matter of a direct correspondence between language (or thought) and a mind-independent world of entities, properties, relations, or, more simply, facts (construed as constellations of the former). For expository purposes, it will be convenient to adopt Putnam's method of capitalizing such words as 'world,' 'object,' 'property,' and 'fact' in order to make clear that when I am talking about the inhabitants of the mind and discourse-independent world it is part of the correspondence view. As we shall see later on, the moral irrealist can help herself to talk about moral properties, relations, and facts, but, of course, it is part of the irrealist's story that there are no moral PROPERTIES, RELATIONS, or FACTS. So, according to the correspondence view, the truth of any sentence is a matter of its corresponding to the WORLD of OBJECTS, PROPERTIES, RELATIONS, and (if one wants to speak this way) FACTS.⁴⁰

Notice that all we are told about the nature of PROPERTIES is that they are mind- and language-independent properties.

Hereafter, I shall use the term 'Timmons-property' to mean what Mark Timmons means by the word 'property' when he puts it in lower case (it is an open question, so far, whether that is the same as what 'property' means in English). Now, I

contend that there is no sense to be made of Timmons' views, unless he is misusing 'property'. Consider what 'property' means in standard English: does it mean PROPERTY, or does it mean Timmons-property?

Case 1: Suppose 'property' in English means PROPERTY. Then the above passage from Timmons gives a circular definition: in effect, it tells us that a PROPERTY is a mind-independent PROPERTY. More importantly, Timmons' central claim that there are moral properties but no moral PROPERTIES is a straightforward contradiction.

One might say that he is not contradicting himself, because when he says this, he is using 'properties' non-standardly (that is, to mean Timmons-properties). But in that case, his account of the meanings of words is incorrect, and he therefore fails in his effort to show that common sense isn't committed to moral realism.

Case 2: Suppose 'property' in English means Timmons-property. Then PROPERTIES, according to the above passage, comprise that subset of Timmons-properties that are mind-independent. Timmons does not define 'mind-independent', but here are two reasonable criteria:

a) \$F\$-ness is mind-independent, provided: If a thing is \$F\$, then it would still be \$F\$ even if we had entirely different attitudes about it.⁴¹

b) \$F\$-ness is mind-independent, provided: It is possible that a thing be \$F\$ even if no one thinks so.

Timmons would affirm the mind-independence of moral properties on both interpretations. He says we are entitled to avow moral fallibility, even of the most radical kinds, so he would accept that moral properties satisfy (b). He also stresses that his view is not relativistic, and that it does not hold that a moral assertion describes the speaker's attitudes. He thinks moral assertions are 'categorical', which implies among other things that the speaker must be able to support them with 'non-subjective reasons, in other words, reasons that do not appeal to matters of individual taste, personal preference, or the like'.⁴² So it seems that on his view, the logic of moral discourse commits one to affirming that moral properties satisfy (a). Moreover, this is obvious in any case—it is false that 'slavery would not be wrong, if only we slavery-haters mellowed out a little'.⁴³ So it seems that, if a PROPERTY is just a Timmons-property that does not depend on the mind, then Timmons does indeed take there to be moral PROPERTIES. So he is a moral realist.

Now here is what I think is really going on. Of course Timmons isn't a realist. We understand what Timmons is saying, and we understand it to differ from realism, because 'property' in English means PROPERTY; capitalization adds nothing. We understand Timmons' real view to be that (a) there aren't actually any moral properties, but (b) we can still go around using the word 'property', because it's just one of the rules of moral language that you get to use that word even when there aren't any properties around for it to refer to. The reason why Timmons' view is initially hard to grasp and why it appears so late on the scene, is that the rules of philosophical discourse preclude his describing his view in the most perspicuous way. For if he describes his view as I have, its falsity is apparent: it is not the case that in English, you are allowed to say there are properties of some kind when in fact there aren't any.

If Timmons' semantic views were correct, then we could not understand what realism was. His views entitle him—indeed, require him—to affirm all the sentences characteristic of realism: 'There are moral facts', 'Moral facts are independent of our attitudes', 'Moral statements are true when they correspond to the facts',^[44] and so on. If the block-letter terms were not the ordinary English terms—if they were technical terms, as Timmons must represent them as being—then he would have to explain in English what they mean. But none of his attempted explanations would succeed, because we would interpret all the expressions in those explanations minimalistically, so we would take the realist to be committed to no more than Timmons is committed to. When Timmons tried to tell us that realists believe in a MIND-INDEPENDENT WORLD, we would just think he meant a Timmons-mind-independent-world.

Recent versions of non-cognitivism are a kind of meta-linguistic trickery. They seek to embrace common sense in words but reject it in substance. No one but Berkeley was fooled when Berkeley claimed to be an anti-skeptic; nor should we be fooled by today's non-cognitivists.

2.8 The introspective evidence

I have characterized cognitivism and non-cognitivism as theses about the meanings of evaluative statements. Assuming that evaluative statements express value judgments, the semantic claims also imply psychological claims: Cognitivism implies that what we call value judgments are a species of belief, as the term 'judgment' suggests; non-cognitivism implies that they are instead a species of emotion, desire, or other non-cognitive state. We might expect to be able to resolve this

dispute by introspection. When we make value judgments, is the experience more like typical emotions, desires, or other non-cognitive attitudes; or is it more like typical experiences of forming beliefs?

The emotivist might cite the fact that we typically experience emotions when we make moral judgments. Indeed, disputes over such matters as the morality of abortion and the justice of capital punishment are among the most emotional of all disputes. This is some evidence for non-cognitivism; after all, if there generally were no emotions associated with moral judgments, I would certainly claim this as evidence for cognitivism. But before endorsing non-cognitivism on this ground, let us seek a more nuanced view of what the two theories predict about our psychology.

First, although non-cognitivism more clearly predicts an association between moral judgment and emotion than cognitivism does, this association is hardly inexplicable on a cognitivist view. It shows only that we care about moral issues, in some cases very strongly. There are many other issues that we also care about strongly. Debates over whether God exists, who shot JFK, and what a particular disputed clause in a contract means, can all call forth strong emotions. But this exerts little if any pressure towards regarding statements about those subjects as mere emotional expressions devoid of factual content.

Second, our emotions about some things are stronger than about other things. If the core explanation for the correlation between moral judgment and emotion is that 'good' and 'bad' are just emotion-expressing terms, then what should we predict about how the strength of moral emotions should vary? It seems that the strength of the emotions associated with moral judgments should be roughly proportional to the perceived level of goodness or badness of a thing and to the confidence of one's moral judgment: the worse one takes a bad thing to be, and the more confident one is in one's judgment, the more strongly one should feel about it. Is this true? In some cases, yes: I feel more strongly about the wrongness of murder than I do about the wrongness of jaywalking. In other cases, no: I feel more strongly about someone's stealing my stereo (and this is a moral emotion, indignation) than I do about Emperor Nero's execution of Octavia in 62 A.D. Truth be told, I have little if any feeling about Octavia's murder. But Nero's execution of Octavia was far more wrong than the stealing of my stereo, and I am quite certain of that. This suggests that the emotions often associated with moral beliefs have more to do with such things as our concern for our own interests, the social groups we like to associate with, and the self-images we like to maintain, than they do with an emotive meaning built into the term 'wrong'. Though I cannot catalog all the things that influence our emotions, it is unclear why the sorts of factors that explain why people are emotional about other matters cannot also account for our emotions about moral issues.

Third, we can distinguish between the degree of confidence of a moral judgment and the degree of rightness or wrongness that the moral judgment attributes. Imagine someone who thinks abortion is likely wrong but lacks confidence in that judgment. He sees that there are plausible arguments on the other side. However, he says, if abortion is wrong, it is very seriously wrong, since the only way it would be wrong would be by being akin to murder. Here, the perceived degree of wrongness is very high, but the confidence in the judgment is very low. On the emotivist theory, what could these two dimensions correspond to? On the emotivist theory, one simply experiences a feeling about abortion. This feeling could come in degrees of intensity—people can feel more or less strongly about something. But how could this give rise to our sense that there are two things here that come in degrees and that may have radically different degrees in a given case?

Fourth, many moral judgments are made dispassionately, with no evident emotion. Now, in the case of Nero's killing of Octavia, the emotivist might say that I have a general sentiment of disapproval towards murder, and this induces me to say Octavia's murder was wrong even though I have no feeling about Octavia in particular. But this leads to a puzzling (for emotivists) feature of moral emotions: the most fundamental of our moral principles tend to elicit the least feeling. Consider the following candidate fundamental moral principles:

Enjoyment is good.

One should treat others as one would wish to be treated.

The right action is always the action that produces the most good.

A being has a right to something only if it can desire that thing.

Whether or not these principles are correct, they are examples of the sort of thing that philosophers take, plausibly enough, to be fundamental moral principles. But few people feel particularly emotional about these sorts of statements. They tend to be far more emotional about certain derivative moral judgments: that Steve mistreated Alice, that the Simpson verdict was unjust, or that the boss' demands are unreasonable. This is unsurprising to the cognitivist, as it is just an instance of a general pattern: people are more emotional about concrete matters affecting their own interests and those of

people they know than they are about abstract, theoretical questions. But for the emotivist the phenomenon is unexpected: if moral judgments are essentially emotional, then shouldn't the most fundamental moral judgments exhibit the clearest, purest moral emotions?

Notice an important asymmetry here. Cognitivism does not require that moral judgments are always dispassionate; we know that people can be influenced by emotions in forming their beliefs, just as their beliefs can in turn influence their emotions; there is nothing special about morality in this regard. But emotivism requires that moral judgments are always emotional—because that, according to the theory, is just what moral 'judgments' are. Introspection shows this to be false.

Fifth, the procedures whereby we at least sometimes form moral judgments—notably, in moral philosophy—are far more like the way we form beliefs than like the way we typically acquire emotions. I want to know whether abortion is really wrong. I go to the moral philosophy literature on the subject. I read Thomson's thought experiment about the violinist. I read Marquis' argument about the rights of potential persons. I read Tooley's argument about the criteria for having a right to life.⁴⁵ In all of this, my mental activity is indistinguishable from what it would normally be when I am trying to figure something out. I start in a state of suspended judgment. I weigh reasons for and against a conclusion. Of course, nothing inherently precludes this sort of activity from resulting in one's feeling some emotion. But this is not normally how emotion works. I don't normally go to the professional literature for the purpose of having an emotion about something. I don't normally carefully weigh arguments, looking out for logical errors, for the purpose of forming an emotion. I don't normally try to decide what emotions to have, still less worry about 'getting it right'. None of those things is characteristic of how I would come to be afraid of heights, or love my girlfriend, or be annoyed by Barney the dinosaur.

As in the case of the linguistic evidence, everything here seems to be exactly as it would be if a cognitive interpretation of moral language were correct. I conclude that ethical terms are cognitively meaningful, just as 'blue' and 'square' are, and that ethical statements are claims about an aspect of reality, just as 'The sky is blue' and 'Dogs are square' are.