

7 Practical Reasons

Part 1: Humean Argument and Rationalist Motivation

Many anti-realists have charged that if moral realism were true, there would be no reason for acting morally. In this chapter, I argue instead that only the intuitionist has a plausible account of our reasons for being moral.

7.1 The Humean argument against realism

What reasons do we have for behaving morally? There are at least two interpretations of this question, because there are two senses of 'reason': there are motivating reasons and normative reasons.

Motivating reasons (or 'motives') are the sort of things that explain why people do what they do. One's having a motivating reason for doing $\$A\$$ does not entail that one will do $\$A\$$; it means only that one has at least some tendency or temptation to do $\$A\$$, given the opportunity. For example, my hunger, combined with my belief that the sandwich before me is food, constitutes a motive for me to eat the sandwich.

Normative reasons are the sort of things that explain why people should do certain things, from the standpoint of reason.¹ A normative reason for performing an action is a consideration in favor of the action, though not necessarily a decisive consideration. Like motivating reasons, normative reasons can conflict. What it is rational to do is what one has the most normative reason to do. For example, the fact that I need food, together with the fact that the sandwich before me is food, constitutes a normative reason for me to eat the sandwich.

We typically have normative reasons and motivating reasons for the same things. In the above example, I have both a motivating and a normative reason for eating the sandwich. But often the strength of our motives fails to match the strength of our normative reasons. Suppose I have a choice between a nutritious veggie sandwich and a delicious ice cream sandwich. I may have better reasons, all things considered, for eating the veggie sandwich, but I may find, alas, that I have more motivation to eat the ice cream. The anticipated gustatory pleasure of the ice cream is some justification for choosing it (insofar as pleasure is a good), but this justification may be rationally outweighed by long-term health considerations. These long-term health considerations produce in me some tendency to choose the veggies, but this tendency may be outweighed, motivationally, by the desire for sugary pleasure, and I may thus make the irrational choice.

The following three principles about reasons for action are widely accepted in philosophy; the first two are standard assumptions in economics as well:

i) The Humean Conception of Reasons (hereafter, 'the Humean Conception'): All reasons for action derive from pre-existing desires. In particular,

a) You have a motivating reason to do $\$A\$$ if and only if there is something you desire for its own sake that you believe $\$A\$$ would or might accomplish. The strength of your reason to do $\$A\$$ is proportional to the strength of that desire.

b) You have a normative reason to do $\$A\$$ if and only if there is something you desire for its own sake that you have reason to believe $\$A\$$ would or might accomplish.²

ii) The Belief/Desire Gap: Beliefs don't entail desires: it is possible to believe any given objective fact obtains, while lacking any given desire.

iii) The Magnetism of Values:

a) If one accepts a moral value, one necessarily has a motivating reason for action. For instance, if you take stealing to be wrong, then you must have at least some motivation to avoid stealing.

b) Moral facts, if there are any, necessarily provide normative reasons for action. If stealing is wrong, or at least if one has reason to believe it to be wrong, then one has a normative reason to avoid stealing.

These principles are commonly thought to refute moral realism. How so? There are two versions of the argument:³

First Version:

1. Motives for action require desires. (The Humean Conception.)
2. No belief about an objective fact entails the presence of any desire. (The Belief/Desire Gap.)
3. Therefore, no belief about an objective fact, by itself, entails the presence of any motive for action. (From 1, 2.)
4. Moral attitudes (the attitudes expressed by moral claims) entail the presence of motives for action. (The Magnetism of Values.)
5. Therefore, moral attitudes are not beliefs about objective facts. (From 3, 4.)

Second Version:

6. Normative reasons for action require desires. (The Humean Conception.)
7. No objective fact entails the presence of a desire. (Because objective facts are facts in the external world, apart from our desires and attitudes.)
8. No objective fact, by itself, entails the presence of any normative reason for action. (From 6, 7.)
9. Moral facts, if there are any, entail the presence of normative reasons for action. (The Magnetism of Values.)
10. Therefore, there are no objective moral facts. (From 8, 9.)

In brief, moral claims state reasons for action. Reasons for action, of either kind, require the existence of desires. But no objective fact, or belief about an objective fact, can by itself entail the existence of any desires. Therefore, moral claims do not merely report objective facts; they must contain some subjective or non-cognitive element.

There are three ways for a moral realist to respond: reject the Humean Conception, reject the Belief/Desire Gap, or reject the Magnetism of Values. Each way has been taken by at least some philosophers.⁴ I shall criticize the Humean Conception and propose an alternative conception of reasons.

7.2 The connection between motivating and normative reasons

A natural response to the first version of the Humean argument is to question (4). If I take stealing to be wrong, then perhaps it follows that I ought to be, or that I think I ought to be, motivated to avoid stealing. But it does not obviously follow that I am in fact motivated to avoid stealing. I might be an immoral person, even by my own lights—that is, I might fail to be motivated by what I take to be morally correct. For similar reasons, we might doubt (6) in the second version. If I have a normative reason to help Jimmy with his homework, perhaps it follows that I should desire to help Jimmy, or that I would desire to do so were I fully rational, but it does not obviously follow that I in fact have some desire that would be satisfied by helping Jimmy. There is reason to suspect that the argument trades on an equivocation: the Humean thesis—that reasons for action require desires—is plausible only if we take ‘reasons’ to mean ‘motivating reasons’. But the Magnetism of Values premise—that moral claims state reasons for action—is plausible only if we take ‘reasons’ to mean ‘normative reasons’. That is, it is plausible that our being motivated to do something requires our having a desire appropriately related to the action, but not that our having a normative reason to do something requires such a desire; and it is plausible that our justifiably taking something to be good must give us a normative reason to pursue it, but not that it must actually motivate us to pursue it. But the argument requires both the Humean Conception and the Magnetism of Values to be true for the same sense of ‘reason’.

That is not my ultimate response to the argument, however. I raise this point to show why a proponent of the Humean argument should want to defend some strong connection between motivating and normative reasons, roughly to the effect that the two kinds of reasons always go together. Fortunately for them, they can do so by means of at least two plausible arguments.

First, an argument from the ‘ought implies can’ principle: To say that you should do A implies that you can do A of your own free will. This implies that you can freely choose to do A, which in turn implies that you can be motivated to do A; a person

cannot freely choose to do something while having no motivation at all for doing it. Now, to say that you have a normative reason for doing A implies that you should do A (provided there are no countervailing considerations). It follows that you are capable of being motivated to do A (provided there are no countervailing considerations). But if you had no motive that would appear to you to be served by doing A, then, since you also have no choice about what motives you have, you would be incapable of being motivated to do A. Therefore, the claim that you have a normative reason for doing A implies that you have some available motivation for doing A. In short:

1. If S has a normative reason for doing A, then S should do A, provided there are no countervailing reasons.
2. If \$\$\$ should do \$\$\$, then \$\$\$ can freely do \$A\$.
3. If \$\$\$ can freely do \$A\$, then \$\$\$ can be motivated to do \$A\$.
4. Therefore, if \$\$\$ has a normative reason for doing \$A\$, then \$\$\$ can be motivated to do \$A\$, provided there are no countervailing reasons. (From 1, 2, 3.)
5. If \$\$\$ has no available motivating reason to do \$A\$, then \$\$\$ can not be motivated to do \$A\$, even if there are no countervailing reasons against \$A\$.
6. Therefore, if \$\$\$ has a normative reason for doing \$A\$, then \$\$\$ has an available motivating reason for doing \$A\$. (From 4, 5.)

To illustrate, suppose someone tells me that I have a normative reason to donate \$500 to Oxfam. Suppose, however, that I am so lacking in concern for my fellow human beings that I have no desire to help the poor, nor any other desire that would be satisfied by donating this money. Suppose that after thorough and correct deliberation, I realize these facts. Furthermore, I cannot help any of this; I cannot help the fact that I don't care about the poor, nor the fact that I know that donating \$500 to Oxfam satisfies none of my desires. In this situation, I cannot be motivated to donate the money. Therefore, I simply cannot donate the money of my own free will. And therefore, it is false to say that I should donate the money. Conversely, for it to be true that I should donate the money, one of the above assumptions would have to be false; I would have to have some desire that I thought donating the money might satisfy.

That is an argument a Humean might make, to draw the connection between motivating and normative reasons, and thereby to bolster the contention that 'S should do A' implies that S has some sort of desire.

A second argument for this connection I call the Argument from Interpretation. Our normal way of interpreting other people's behavior (including their statements) relies on an assumption of rationality: when you are trying to figure out what another person believes, desires, and so on, you look at his behavior and attribute to him the beliefs, desires, and so on that would make that behavior make the most sense. For example, if someone eats a lot of chocolate, I infer that she likes chocolate and thus is undertaking the correct means of satisfying her preferences. If I instead assumed that she was irrational, I might hypothesize that she hates chocolate and that is why she eats it. A similar point applies to attributing beliefs to others: I assume my chocolate-eating friend believes that she is eating chocolate, because that is what it would make sense for her to believe, given her actual circumstances. If I took her to be irrational,

I might instead hypothesize that she thinks she is swimming the English Channel, and she likes swimming the Channel, so that is why she eats the chocolate. The hypotheses that assume irrationality might 'explain' the same behavior, but we rightly regard them as *prima facie* ridiculous. The point is that we need to assume others are at least generally rational in order to select an appropriate account of their behavior from among the indefinitely many possible accounts.

Something like this principle is widely accepted under the name of 'the Principle of Charity'. The principle is not that all actions or beliefs are rational, but only that we should interpret others to be by and large rational, if possible. Sometimes, people's statements and behavior provide evidence that they are irrational in certain respects. The Principle of Charity does not deny this, but it holds that rationality is a default assumption, and that we should not attribute more irrationality than necessary. The point is not merely one about how we in fact interpret others, but about what is the correct way of interpreting others. Notice also that the principle is not an empirical one: it is not that we have, by observing people for a long time, discovered that they are by and large rational. Rather, the assumption of general rationality is one that we bring to our observations, as a constraint on their interpretation. We do not directly observe others' beliefs and desires, so we would have no way of finding out whether people were by and large rational, without first having some way of inferring their mental states from their behavior. That sort of inference is just what 'interpretation' is. Since the assumption of rationality is required for interpreting behavior, the assumption of rationality cannot be tested by observing people's behavior. At

most, what could be discovered empirically is that people exhibit behavior which is capable of a rational interpretation, but we cannot determine empirically that interpretations of people as rational are pro tanto better than competing interpretations of people as irrational.

The Principle of Charity is almost universally accepted. But what is its basis? It seems that the principle could be correct only if we were justified in believing, in general, that people are by and large rational. But for the reasons just discussed, this belief cannot be justified by observation. So it must be *a priori*. How can that be? Here is a natural answer: because it is a necessary truth that whenever one has a normative reason for action, one has a motivating reason for action. Suppose the contrary: suppose, that is, that one can have a normative reason to do $\$A\$$ but no motive for doing $\$A\$$. If this is possible, then it seems that it would also be possible for a person to lack motives, in general, for doing the things he had normative reasons for doing. Such a person would be generally irrational. But if it is possible for a person to be generally irrational, then how could we know *a priori* that people are in fact by and large rational?

The preceding argument may be weakest at its first stage, where I construed the Principle of Charity to imply that we should interpret people as being by and large rational. Perhaps instead we need only interpret people as being rational *in certain respects*, which leaves open the possibility that we might discover other broad, systematic respects in which people are not rational. However, I shall not pursue this line. I find it plausible that normative reasons have the kind of 'magnetism' the Humean argument posits. The idea is not that all beings are moved by all normative reasons. The idea is that if a being can and does *recognize* a normative reason, then the being must have at least some tendency to be moved by that reason. But there may be many beings who fail to recognize many normative reasons.

So far, I have sought to state the Humean case as strongly as possible: it seems that evaluative claims imply the existence of normative reasons for action; this, in turn, implies or strongly suggests the existence of motivating reasons for action; and that implies the existence of desires that are appropriately related to the actions in question. But no objective fact can imply the existence of a desire; therefore, evaluative claims are not objective, factual claims. Where I challenge this argument is in its assumption that all motivation stems from desire.

7.3 A rationalist conception of motivation

There are at least four kinds of motivations we are subject to:

i) *Appetites*: Appetites include such things as hunger, thirst, and lust. They are simple, often instinctive desires, whose satisfaction is normally attended by physical (that is, sensory⁸) pleasure, and/or whose frustration is attended by physical pain or discomfort.

ii) *Emotions*: Emotional desires, such as those occasioned by anger, fear, and love, differ from appetites in that their satisfaction or frustration is normally associated with a different (emotional) kind of enjoyment or suffering, rather than sensory pleasure or pain. In addition, emotional desires require a more sophisticated kind of cognition than appetites. To be angry at someone involves perceiving something he has done as being bad in a certain way; to fear something involves perceiving it as dangerous; to love someone involves perceiving his existence and one's relationship to him as valuable; and so on. No such sophisticated perceptions are required for the appetites. Correspondingly, emotions can have much more sophisticated objects; one may be angry about a fact that requires quite abstract concepts to be aware of—say, the fact that a student plagiarized an assignment. Our emotions commonly have longer-term effects as well—said anger may result in one's giving the student a low grade, months later. Appetites have simpler objects and shorter-term effects.

iii) *Prudence*: Prudence is the motivation to pursue (or avoid) a thing because one recognizes it as furthering (or setting back) one's own interests. Prudence differs from, and often opposes, appetitive and emotional desires. To return to an earlier example, my appetite for ice cream may conflict with my prudential motive to maintain good health. Likewise, my anger at someone may tempt me to hit him, while prudence bids me to restrain myself.

iv) *Impartial Reasons*: Sometimes we are motivated to act in particular ways by our recognition of what is good, fair, honest, and so on. These reasons are impartial, in the sense that they don't appeal to our own interests or to our personal preferences. This sort of motive may conflict with any of the first three; for instance, a man may desire to cheat on his wife, but believe that this would be unfaithful (and hence bad). Or a judge may consider taking a bribe, which would promote his own interests, but believe that this would be dishonest (and hence bad).

These seem to be the main kinds of motives. There are several interesting questions about them. First, what is the object of an appetitive desire—that is, when one has an appetite, what is it that one wants? One traditional view is that it is

physical pleasure or the avoidance of physical pain, and this has led some thinkers to the generalization that all desire boils down to the desire for pleasure and the absence of pain. This, however, is a mistake. One may of course desire pleasure, but this is not what typical appetites consist in, even though pleasure normally attends the satisfaction of those appetites. When you are hungry, what you want is food. You may also want gustatory pleasure, but that desire is not what your hunger consists in; hunger consists in the desire for food. To illustrate, suppose you are hungry. Some mad scientists stimulate your brain so that you have an incredibly realistic hallucination of eating a salmon fillet, and so that at the same time you stop feeling hungry. In this case, I would say that your desire (the hunger) went away without being satisfied—you did not get what you wanted, although you may have been tricked into thinking you did. Our desires are typically focused on the world, not on our minds.

Similarly, what are the objects of emotional desires? If one is tempted to view typical appetites as desires for physical pleasure, one might be similarly tempted to treat emotional desires as desires for emotional satisfaction. Thus, one might view a man's love for his children as giving rise chiefly to a desire for the emotional satisfaction of seeing them well off. But in fact, what love gives rise to is the desire that loved ones be better off. That is why, for example, taking steps to ensure that one's children are provided for after one's death is a manifestation of love.

To illustrate, imagine that you are offered a chance to sell your children into slavery for \$10,000. The children will be miserable for the rest of their lives. However, as part of the deal, you will receive a hypnosis session that will cause you to forget that the children were sold into slavery and will, furthermore, induce in you a permanent delusion that the children are extremely happy and successful. You will thereby enjoy the satisfaction of thinking your children to be well off. Assuming that you love your children very much, would you take the deal? If you answered 'no', then your chief desire is not for emotional satisfaction.[11]

Where does the desire for pleasure fit into my taxonomy of motivations? I think that it fits under prudence, in the admittedly broad sense in which I have defined 'prudence'. The desire to secure pleasures for oneself is a desire to promote one's interests in a certain way. We may also just have an appetite for pleasure.

This brings us to another simplistic but common psychological theory: the theory that all actions are motivated by concern for self-interest. This theory falls to much the same objection as the previous one. We are moved by a variety of different desires—desires for food, sex, power, wealth, love, fame, and so on—most of which are for things logically independent of the advancement of our interests. No *recherché* or even particularly noble examples are required to refute psychological egoism. Consider a desire, prompted by anger, to hit someone: one need not take such an action to be in any way in one's interests in order to feel the desire, and one may well take it to be very much against one's interests.[12]

These remarks undermine two particular Humean theories of reasons: the view that all actions are motivated by the desire to secure pleasure and avoid pain, and the view that all actions are motivated by the desire to promote one's own interests. But they leave intact the more general Humean framework, according to which all actions are motivated by desires of some kind.

Where does morality fit into the taxonomy? I take moral principles to be a species of impartial reasons. Thus, considerations of what is good, fair, honest, and so on are both moral reasons and impartial reasons for action. There may be other impartial reasons that are not moral (perhaps the fact that a certain procedure is the most efficient way of producing pencils is an impartial reason, but not a moral reason, for one to cause the pencil factory to make pencils in that way), but I won't discuss them here.

We always have an impartial reason for doing what we know to be morally correct. But this is not to say that that reason is the only, or even the primary reason we have for acting morally. We also have many non-moral reasons for being moral. Consider, for example, a person who does what he believes is right, not merely because it is right, but because he believes God will punish wrongdoers; or a person who refrains from stealing, not merely because he recognizes theft as wrong, but because stealing makes him feel bad. These would be, respectively, prudential and emotional motives for doing what is right.

How does the notion of 'desire' map onto my taxonomy of motives? In ordinary language, appetites and emotional motives are commonly called 'desires', while prudential and impartial motives are not. Imagine a student who knows she has an exam tomorrow morning, for which she has not studied. A friend invites her to a party. The dialogue might go like this:

Friend: Hey, you wanna go to a party tonight?

Student: I really want to go, but I should study tonight.

Friend: So you're saying you'd rather study than party?

Student: Of course I don't want to study; I'd much rather be partying. But I really need to study.

The natural interpretation is that the student has described her motivations correctly: she has an emotional desire to party, no desire to study, but recognizes—that is, believes—that it would be very much in her interests to study (she 'needs' to). The latter, prudential motivation may overrule her desire.

Or imagine a judge who has been offered a bribe to acquit a criminal. The dialogue might go like this:

Judge: I won't take your bribe.

Criminal: What's the matter, don't you want \$100,000?

Judge: Of course I want \$100,000. But acquitting you would be wrong.

Criminal: So you're saying that you want to send me to prison more than you want to have \$100,000? You must really hate me.

Judge: It doesn't matter how I feel about you. You committed a crime, and criminals should be punished, and it's my job to see to it.

Criminal: Look, I bet my money could buy you a lot more enjoyment than whatever satisfaction you'd get from seeing me punished.

Judge: You just don't get it.

The natural interpretation is that the judge recognizes that taking the bribe would best promote his own interests, satisfy his desires, and increase his enjoyment, but he believes it would be wrong. His moral belief overrules his personal interests.

What would the Humean say? That in the first example, the student wants to get a good grade on the test, more than she wants to go partying, and that in the second example, the judge wants the criminal to be punished more than he wants the \$100,000. Is this true? Not in any non-trivial sense. 'Desire', in English, seems often to refer to appetites and emotions, as the above examples illustrate. In that sense, the Humean's claim is false. The student may well feel more desire to go partying than to get a good grade on the test, but still recognize that it would be more in her interests to study. The judge may feel more desire to take the \$100,000 than to punish the criminal, but recognize that justice requires him to punish the criminal. Introspectively, there is simply no evidence for the Humean claim.

In another sense, however, a 'desire' is simply any motivation, of whatever kind.[13] In this sense, it is trivial that in my examples, the student desires to study and the judge desires to refuse the bribe. On this interpretation, the core Humean thesis that motives for action require desires amounts to no more than the following tautology, with which even the staunchest rationalist will agree: Motives for action require motives for action. But surely no interesting conclusion can be inferred from this. In particular, on the present conception of desire, it can no longer be assumed that beliefs are not desires; if a belief motivates action, then it would ipso facto be a 'desire' in the present sense. And we should simply have to distinguish two kinds of 'desires': the kind that are constituted by beliefs, and the kind that are constituted by appetites or emotions.

I take it that the Humean does not want to trivialize his thesis in this way; most Humeans will say that we have some independent conception of 'desire', as an introspectively identifiable type of mental state, on which it is not an open question whether a desire can be a belief.[14] But on introspecting, the only such mental states I can find that might be called 'desires' are the appetites and emotions. And it is not debatable whether I always act on my currently strongest appetites and emotions—I know that, thankfully, I do not.

7.4 Why believe the Humean conception?

7.4.1 The intuitive appeal of the Humean conception

Why have so many philosophers embraced the Humean Conception of Reasons? Most seem simply to consider it intuitively obvious—it is just obvious that reasons for action require goals, and it is just obvious that only our desires, and not mere

beliefs about the world outside of us, can provide us with goals.¹⁵ Let us first discuss this appeal to intuition, and then turn to some arguments for the Humean Conception.

I am at least as willing as the next person to consider intuitions. But I remain unpersuaded by the widespread intuition supporting the Humean Conception. This intuition is a *general* one; that is, its content is a universal generalization ('All reasons have such-and-such form' or 'No belief can be inherently motivating'). I think that intuitions of this kind can often be easily undermined by particular examples of things that defy the generalization. For example, consider the principle, 'Two solid objects cannot occupy the same space at the same time'. On its face, this seems obviously true. But now consider this case: you have a lump of clay, which we can call 'Lumpy'. You proceed to mold the clay in such a way as to form a statue. Call the statue 'Venus'. It seems that Lumpy still exists; it has simply been molded into a different shape. Furthermore, Lumpy and Venus now occupy the same space. Both are solid objects. And they are distinct objects, since Lumpy existed before Venus did. Therefore, we have two solid objects occupying the same space at the same time. Isn't this a deep philosophical paradox?

No, it isn't. We should respond to the puzzle by saying: 'Okay, if you're going to count cases like that, then two solid objects can occupy the same space at the same time. When I said this was impossible, I had in mind pairs of objects like: my hand and this table. If I had been thinking of "pairs of objects" like the statue and the mass of clay it is composed of, I would never have said that two solid objects couldn't occupy the same space, because obviously those things can'. Thus, I think that the general intuition we began with—even though it was quite a firm and widely-shared intuition—has no probative force against the judgment that Lumpy and Venus are in the same place, or that they are two solid objects. The reason is that the general intuition depended on our having certain kinds of cases, and not others, in mind in the first place.

I am not claiming that this is true of *all* general intuitions, but it is true of many of them. And in particular, I think it is true of the intuition supporting the Humean Conception. When someone puts forward the general principle, 'Beliefs cannot by themselves be motivating', I think that we can only judge the plausibility of that generalization by running through different kinds of beliefs in our minds. To be sure, the principle seems plausible when we think about *most* beliefs—for instance, the belief that my house is on fire, the belief that the weasels are mammals, the belief that 13 is prime, and so on. In fact, it seems plausible when you consider any belief about any purely natural, external fact. But now consider another class of beliefs: beliefs about *value*. For instance, the belief that stealing is *morally wrong*, the belief that happiness is *good*, and so on. If you're counting *this* sort of belief, then the principle does not seem at all obvious; indeed, it seems very implausible. I think the principle that only desires can motivate seems obvious only when we are ignoring certain cases—just as the principle that two solid objects can't occupy the same space only seems obvious when we are ignoring certain cases.

Consider why it seems plausible that a belief about a natural fact cannot be inherently motivating. One reason is that it is unclear what would determine *which* course of action the natural fact is supposed to favor. Thus, suppose you see someone sitting in a train station. You inform him that the train for New Brunswick is about to leave. Can this fact, by itself, motivate him to act in some way—for instance, to get up and board the train? One reason for denying that it can is that there doesn't seem to be any particular course of action that uniquely bears any special relation to the fact. Granted, the fact that the train is about to leave implies that getting up and boarding the train is required *in order to get to New Brunswick on that train*. On the other hand, refraining from getting up and boarding the train is required *in order to keep out of New Brunswick*. The mere fact that the train is about to leave does not tell us whether getting to New Brunswick is good or bad, so it doesn't pick out any particular course of action as the one to take.

But evaluative facts are obviously different. They *do* tell us what is good or bad, and there is no puzzle about which course of action an evaluative judgment supports. The judgment that one should keep promises obviously bears a special relationship to promise-keeping that it does not bear to other kinds of actions.

Now consider Hume's main argument against the Rationalist

Conception: Hume says that the function of reason is only to determine what is true and false. But actions cannot be true or false. Therefore, actions cannot be evaluated by reason, and therefore they cannot be rational or irrational.¹⁶

That is an awful argument. Consider a parallel argument: Reason can only judge truth and falsity. Numbers cannot be true or false. Therefore, reason cannot make any judgments about numbers—for instance, reason cannot tell us that seven is prime. Presumably Hume would say that this is a misapplication of his principle. Presumably, he would say that although numbers themselves cannot be true or false, propositions about them, including the proposition that seven is prime, can be true or false; and that is why the judgment that seven is prime can fall under the purview of reason. But if that's right, then

a parallel response applies to Hume's argument: of course actions themselves cannot be true or false, but propositions about them, including the proposition that some particular action is right, can be; and that is why the judgment that an action is right can fall under the purview of reason.

*7.4.2 Smith's argument

Michael Smith provides an argument similar in spirit to Hume's, but stronger. Smith's argument begins with a technical notion, that of 'direction of fit'. Here is a story to illustrate that notion. There are two men in a supermarket. One has a shopping list and is going through the supermarket collecting items from the list. The other is a detective who is tracking the movements of the first man. The detective makes a list of the things the shopper buys. By the time the shopper gets to the checkout stand, both men have identical lists in their hands. Furthermore, both of them want the items on their list to match what the shopper has in his cart. But there is a difference: the shopper is supposed to make the world match his list; the detective is supposed to make his list match the world. Thus, if the detective should notice that his list says 'Twinkies' but the shopper has no Twinkies in the cart, he can fix the problem by simply crossing out the word 'Twinkies'. But if the shopper should notice that his list says 'Twinkies' but he has no Twinkies in the cart, he cannot fix the problem by crossing out 'Twinkies'; rather, he has to go back and put some Twinkies in the cart. We might express this difference by saying that the shopper's list has world-to-list direction of fit; the detective's list has list-to-world direction of fit.¹⁷

Now, there is something analogous to be said about the difference between beliefs and desires: Beliefs, we might say, have 'mind-to-world direction of fit', meaning that beliefs aim to match reality. If your belief fails to correspond to reality (that is, it is false), then the belief is thereby defective and should be changed. Desires, in contrast, have 'world-to-mind direction of fit', meaning that we aim to make reality match our desires. If a desire fails to correspond to reality (that is, it is unsatisfied), the desire is not thereby defective; rather, one has reason to try to change the world.^[18]

With this background, we can state Smith's argument against the Rationalist Conception of motivating reasons:

1. Having a motivating reason requires having a goal.
2. Having a goal is being in a state with world-to-mind direction of fit.
3. Beliefs have mind-to-world direction of fit, not world-to-mind.
4. Therefore, beliefs cannot be motivating reasons.^[19]

(4) conflicts with the rationalist conception, since rationalists say that a moral belief can constitute a motivating reason.

Where might a rationalist challenge this argument? I think the crucial premise is (2).^[20] First, let us clarify (2) by replacing the metaphorical expression 'world-to-mind direction of fit' with its literal meaning. I take it that the literal meaning of (2) is roughly this:

Having a goal is being in a state such that, if one notices that the world does not match that state, then one will have a tendency to try to change the world to remove the discrepancy.^[21]

But this remains ambiguous. What does it mean for a mental state to 'match' or 'fit' the world? This might mean either that the content of the mental state is realized, or that the goal that the mental state represents is realized. This gives rise to the following two readings of (2):

2_{C} Having a goal is being in a state such that, if one notices that the content of that state is not realized, then one will have a tendency to try to change the world to bring about the content of that state.

2_{G} Having a goal is being in a state such that, if one notices that the goal represented by that state is not realized, then one will have a tendency to try to change the world to bring about that goal.

To illustrate the distinction, suppose I have the goal of increasing my happiness. According to the Rationalist Conception, my having this goal might consist simply in my believing that my happiness is good.²² Now, the content of this belief—the thing I believe to be the case—is: my happiness is good. But the goal that the mental state represents is simply my happiness. This is just part of the rationalist view—the rationalist view is that one's goal is the thing that one believes to be good. Thus, (2c) implies: if I notice that happiness is not in fact good, then I will have a tendency to try to bring it about

that happiness is good. (2G), on the other hand, implies that, if I notice that I am not in fact happy, then I will try to bring it about that I am happy.

Obviously, if (2) means (2c), then the rationalist should reject the premise. If I somehow discover that happiness is not good, I will not try to bring it about that happiness is good (how could I do that? I cannot change the facts of ethics). On the other hand, if (2) means (2G), we should accept it, but then we should reject premise (3). For, to avoid equivocation, (3) will now have to be read:

3G. If one notices that the goal represented by a belief is not realized, then one will have a tendency to modify one's mind, not to try to change the world to bring about the goal.

This is false. If I notice that my happiness (the goal represented by my 'my happiness is good' belief) is not realized, then I will have a tendency to try to change the world to bring about that goal, that is, to bring it about that I am happy.

Why does Smith not notice the ambiguity in (2)? Perhaps because, on his view of motivation, there is no such ambiguity. On the Humean view, when one has a goal, the content of one's mental state is the same thing as the goal that it represents. For instance, when I want to take over the world, the content of the desire is 'that I take over the world', and the goal it represents is also that I take over the world. So, on the Humean view, there is no distinction to be drawn between (2c) and (2G). But there is a distinction on the rationalist view. When one has a moral belief, if the goal that is represented is G, then the content of the belief is not G, but rather 'G is good'.

Note that I am not, in this section, trying to refute the Humean Conception. I am only trying to show that Smith does not succeed in refuting the Rationalist Conception. For Smith to succeed, he cannot assume the Humean Conception to start with. Thus, he cannot assume that (2_{C}) and (2_{G}) are equivalent. So he must tell us which of those he means. If we are taking the Rationalist Conception seriously, there is no reason why we would accept both (2) and (3) of Smith's argument: depending on how the notion of a mental state's 'matching' reality is read, we will either find (2) obviously implausible (if it is read as (2_{C})) or find (3) obviously implausible (if it is read as (3_{G})). Smith therefore fails to provide a non-question-begging argument against the Rationalist.

So far, I have criticized the motivations for the Humean Conception. In the next two sections, I propose to demonstrate its deep implausibility.